

CS 480

Introduction to Artificial Intelligence

November 8, 2022

Announcements / Reminders

- Follow Week 11 TO DO List
- Written Assignment #03 due on ~~Sunday (11/06/22)~~ **Thursday (11/10)** at 11:00 PM CST
- Programming Assignment #02 due on **Sunday (11/20/22)** at 11:00 PM CST
- Quiz #03 posted due on **Sunday (11/13/22)** at 11:00 PM CST
- Grading TA assignment:
https://docs.google.com/spreadsheets/d/1ExS0bKnGt_fdf4LHa3YS1qRA7-lq4xqXVjfSAPMaGVk/edit?usp=sharing
- UPDATED Final Exam date:
 - **December 1st, 2022 (last week of classes!)**
 - **Ignore the date provided by the Registrar**

Plan for Today

- **Bayesian/Belief Networks**
- **Inference in Bayesian Networks**
- **Decision Networks**

Bayesian (Belief) Network

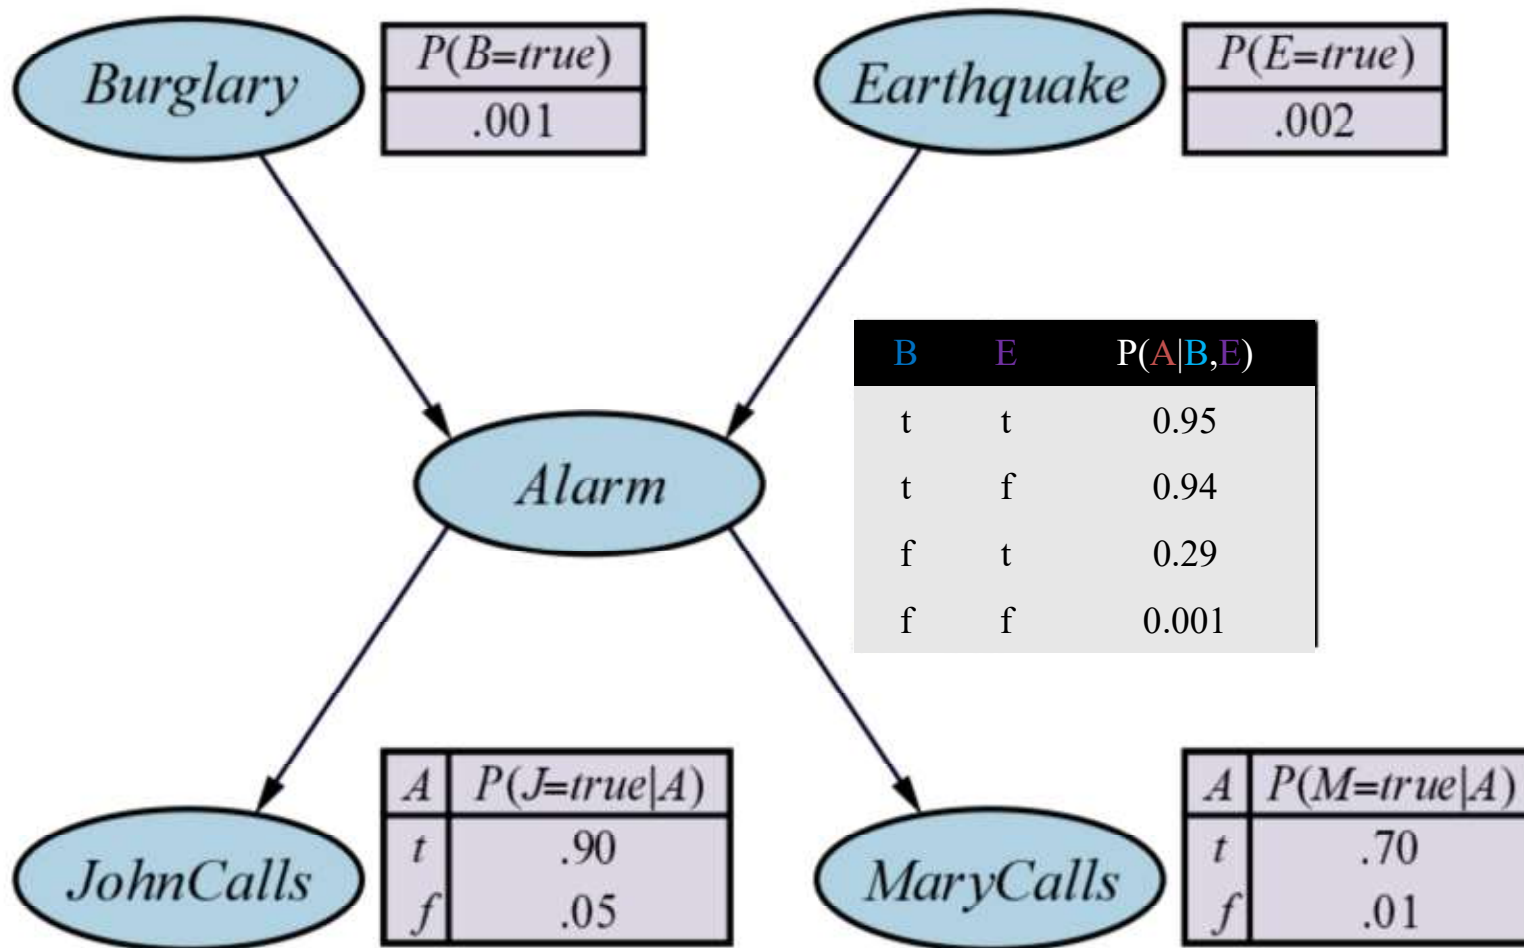
A Bayesian belief network describes the joint probability distribution for a set of variables.

A Bayesian network is **an acyclic, directed graph (DAG)**, where the nodes are random variables (propositions). There is an edge (arc) from each elements of $\text{parents}(X_i)$ into X_i . Associated with the Bayesian network is a set of conditional probability distributions - the conditional probability of each variable given its parents (which includes the prior probabilities of those variables with no parents).

Consists of:

- a graph (DAG) with nodes corresponding to random variables
- a domain for each random variable
- a set of conditional probability distributions $P(X_i \mid \text{parents}(X_i))$


Add Conditional Probability Tables



Inference by Enumeration: Example

Query:

$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$



General Inference Procedure

Given:

- a query involving a single variable X (in our example: **Cavity**),
- a list of **evidence** variables E (in our example: just **Toothache**),
- a list of **observed** values e for E ,
- a list of remaining **unobserved** variables Y (in our example: just **Catch**),

where X , E , and Y together are a **COMPLETE** set of variables for the domain, the probability $P(X \mid E)$ can be evaluated as:

$$P(X \mid e) = \alpha * P(X, e) = \alpha * \sum_y P(X, e, y)$$

where y s are all possible values for Y s, α - normalization constant.

$P(X, e, y)$ is a subset of probabilities from the joint distribution

Inference

Query:

$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Given:


- a query involving a single variable X
- a list of **evidence** variables K ,
- a list of **observed** values k for K ,
- a list of remaining **unobserved** variables Y

the probability $P(X \mid K)$ can be evaluated as:

$$P(X \mid k) = \alpha * P(X, k)$$

$$= \alpha * \sum_y P(X, k, y)$$

where y s are all possible values for Y s, α - normalization constant.



Inference

Query:

$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Given:


- a query involving a single variable X : *Burglary*
- a list of **evidence** variables K : *JohnCalls*, *MaryCalls*
- a list of **observed** values k for K : *johnCalls*, *maryCalls*
- a list of remaining **unobserved** variables Y : *Earthquake*, *Alarm*

the probability $P(X \mid K)$ can be evaluated as:

$$P(X \mid k) = \alpha * \sum_y P(X, k, y)$$

where y s are all possible values for Y s, α - normalization constant.

$P(B)$	$P(\neg B)$	$P(E)$	$P(\neg E)$
0.001	0.999	0.002	0.998



B	E	$P(A B,E)$
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	$P(J A)$	A	$P(M A)$
t	0.90	t	0.70
f	0.05	f	0.01

Inference

Query:

$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$


Given:

- a query involving a single variable X : B
- a list of **evidence** variables K : J, M
- a list of **observed** values k for K : j, m
- a list of remaining **unobserved** variables Y : E, A

the probability $P(X \mid K)$ can be evaluated as:

$$P(X \mid k) = \alpha * \sum_y P(X, k, y)$$

where y s are all possible values for Y s, α - normalization constant.



Inference

Query:

$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$


Given:

- a query involving a single variable B
- a list of **evidence** variables $K: J, M$
- a list of **observed** values k for $K: j, m$
- a list of remaining **unobserved** variables $Y: E, A$

the probability $P(B \mid J, M)$ can be evaluated as:

$$P(B \mid j, m) = \alpha * \sum_e \sum_a P(B, j, m, e, a)$$

where ys are all possible values for Ys , α - normalization constant.



Inference

Query (let's change it a bit for simplicity):

$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Given:

- a query involving a single variable B
- a list of **evidence** variables $K: J, M$
- a list of **observed** values k for $K: j, m$
- a list of remaining **unobserved** variables $Y: E, A$


the query can be evaluated as:

$$P(b \mid j, m) = \alpha * \sum_e \sum_a P(b, j, m, e, a)$$

By Chain rule:

$$\begin{aligned} &P(b, j, m, e, a) \\ &= P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$

P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

Inference

Query (let's change it a bit for simplicity):

$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Given:


- a query involving a single variable B
- a list of **evidence** variables $K: J, M$
- a list of **observed** values k for $K: j, m$
- a list of remaining **unobserved** variables $Y: E, A$

the query can be evaluated as:

$$P(b \mid j, m)$$

$$= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a)$$

P(B)	P($\neg B$)	P(E)	P($\neg E$)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

Inference

Query (let's change it a bit for simplicity):

$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Given:

- a query involving a single variable B
- a list of **evidence** variables $K: J, M$
- a list of **observed** values k for $K: j, m$
- a list of remaining **unobserved** variables $Y: E, A$


the query can be evaluated as:

$$P(b \mid j, m)$$

$$= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a)$$

$$= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a)$$

P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

Inference


Query (let's change it a bit for simplicity):

$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$

P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):

$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$


Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(\neg B)
0.001	0.999

P(E)	P(\neg E)
0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)
t	0.90
f	0.05

A	P(M A)
t	0.70
f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

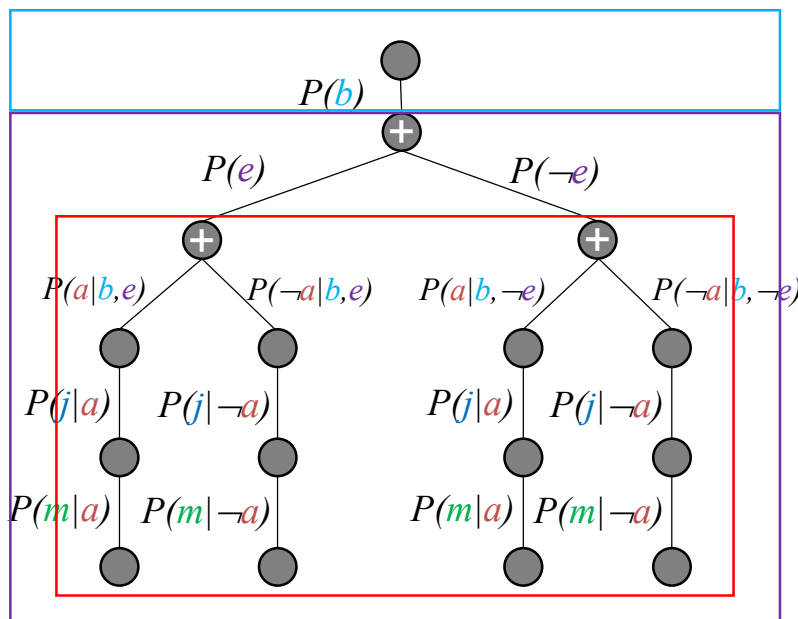
Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01


Inference

Query (let's change it a bit for simplicity):


$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

Query rewritten:

$$\begin{aligned} P(b \mid j, m) \\ &= \alpha * \sum_e \sum_a P(b) * P(e) * P(a \mid b, e) * P(j \mid a) * P(m \mid a) \\ &= \alpha * P(b) * \sum_e P(e) * \sum_a P(a \mid b, e) * P(j \mid a) * P(m \mid a) \end{aligned}$$



P(B)	P(¬B)	P(E)	P(¬E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

Inference

Query (let's change it a bit for simplicity):

$$P(\text{Burglary} = \text{true} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$


We can now calculate:

$$P(b \mid j, m) = \alpha * 0.00059224$$

And through similar process (not shown):

$$P(\neg b \mid j, m) = \alpha * 0.0014919$$

P(B)	P(\neg B)	P(E)	P(\neg E)
0.001	0.999	0.002	0.998



B	E	P(A B,E)
t	t	0.95
t	f	0.94
f	t	0.29
f	f	0.001

A	P(J A)	A	P(M A)
t	0.90	t	0.70
f	0.05	f	0.01

Inference

Query (now we can get joint distribution):


$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

We can now calculate:

$$P(B \mid j, m) = \alpha * < 0.00059224, 0.0014919 >$$

And after normalization:

$$P(B \mid j, m) \approx < 0.284, 0.716 >$$



Inference

Query (now we can get joint distribution):


$$P(\text{Burglary} \mid \text{JohnCalls} = \text{true} \wedge \text{MaryCalls} = \text{true})$$

We can now calculate:

$$P(B \mid j, m) = \alpha * < 0.00059224, 0.0014919 >$$

And after normalization:


$$P(B \mid j, m) \approx < 0.284, 0.716 >$$



Playing Minesweeper with Bayes' Rule


Prior probability / belief:

$$P(X = \text{mine}) = 0.5$$

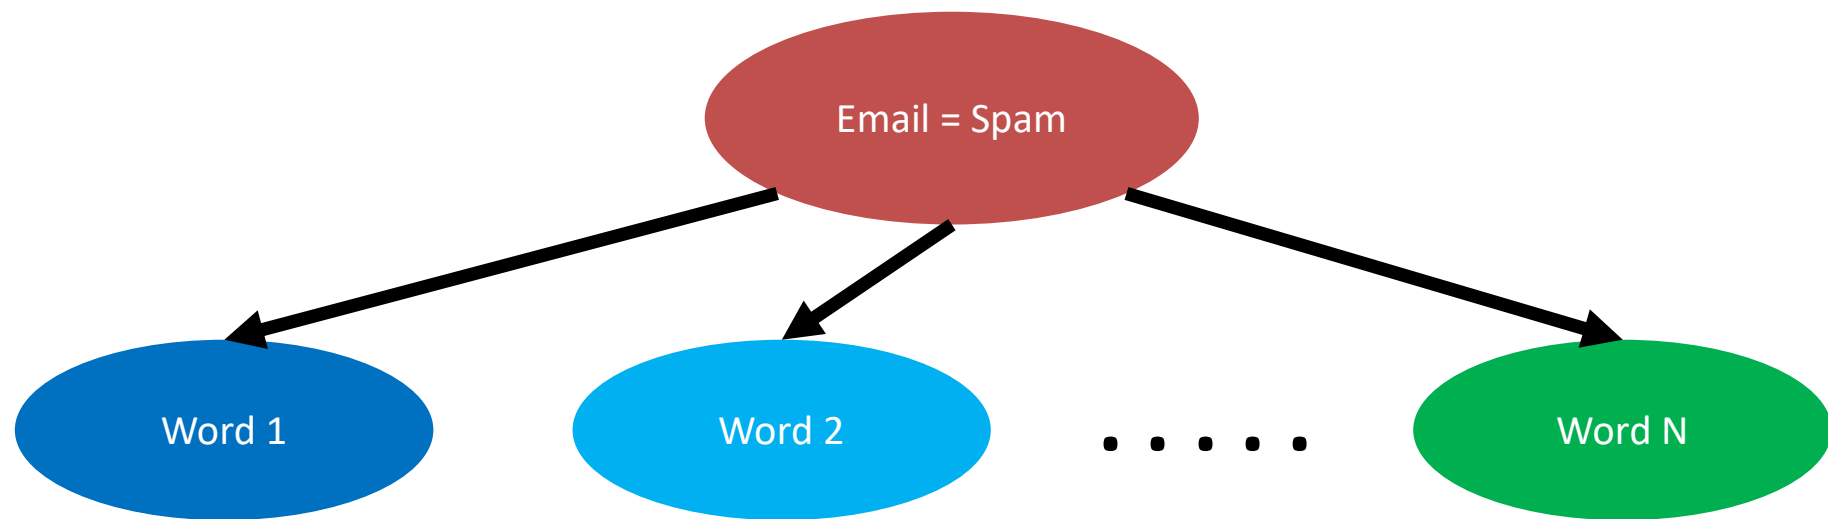


Posterior probability / belief:

$$P(X = \text{mine} \mid \text{evidence}) = 1.0$$



Naive Bayes Spam Filter



$$P(\text{Email} = \text{spam} \mid \text{Word1}) = 0.09$$


$$P(\text{Email} = \text{spam} \mid \text{Word2}) = 0.01$$

...

$$P(\text{Email} = \text{spam} \mid \text{WordN}) = 0.03$$

Conditional Independence


Common **Cause**:



JohnCalls and MaryCalls
are **NOT** independent

JohnCalls and MaryCalls are **CONDITIONALLY**
independent given Alarm


Common **Effect**:




Burglary and Earthquake
are independent

Burglary and Earthquake are **NOT**
CONDITIONALLY independent given Alarm

Conditional Independence



Node X is conditionally independent of its **non-descendants** given its **parents**.



Node X is conditionally independent of ALL other nodes in the network given its **Markov blanket**.

Why do we care?

An unconstrained joint probability distribution with N **binary** variables involves 2^N probabilities. Bayesian network with at most k parents per each node (N) involves $N * 2^k$ probabilities ($k < N$).

Chain Rule

Conditional probabilities can be used to decompose conjunctions using the chain rule. For any propositions

f_1, f_2, \dots, f_n :

$$P(f_1 \wedge f_2 \wedge \dots \wedge f_n) =$$

$$P(f_1) *$$

$$P(f_2 \mid f_1) *$$

$$P(f_3 \mid f_1 \wedge f_2) *$$

...

$$P(f_n \mid f_1 \wedge \dots \wedge f_{n-1}) =$$

$$= \prod_{i=1}^n P(f_i \mid \text{Parents}(f_i)) \quad \leftarrow \text{Enabled by conditional independence}$$

Conditional Independence

Causal Chain:

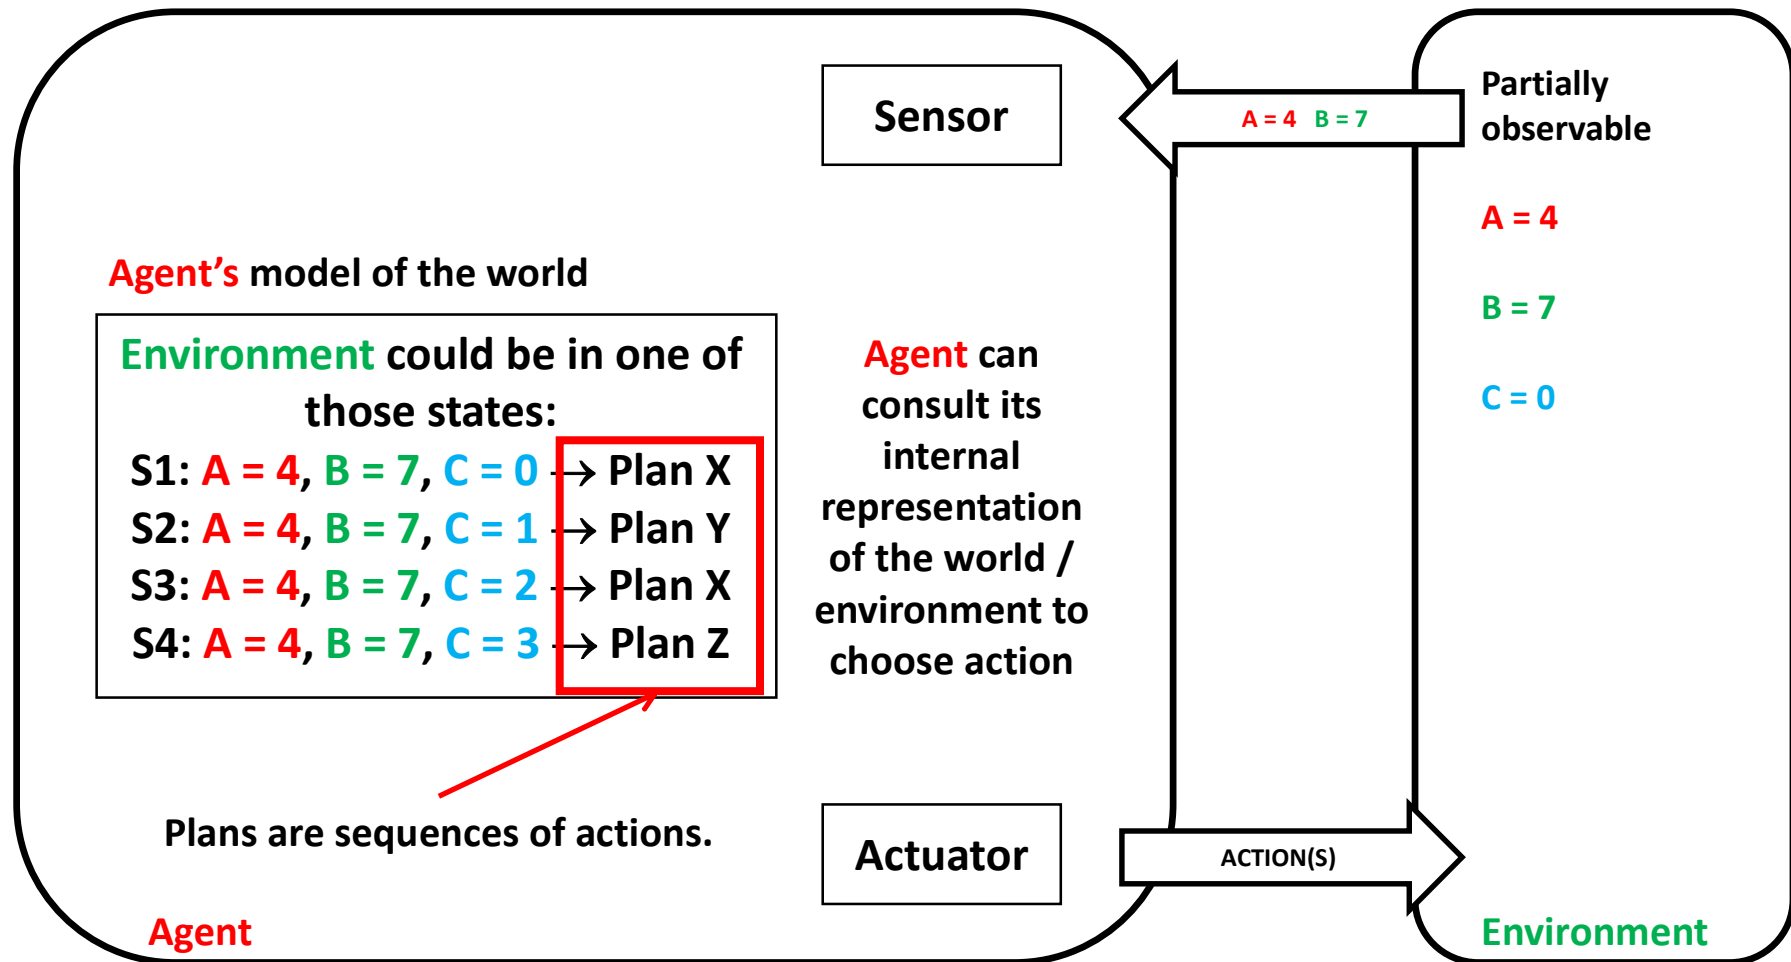


$$P(M | A, B) = \frac{P(A, B, M)}{P(A, B)} = \frac{P(B) * P(A | B) * P(M | A)}{P(B) * P(A | B)} = P(M | A)$$

Burglary and **MaryCalls** are **CONDITIONALLY** independent given **Alarm**.

If **Alarm** is given, what “happened before” does not directly influence **MaryCalls**.

Agents and Belief State



Assume: $D_c = \{0,1,2,3\}$

Decision Theory

- **Decisions**: every plan (**actions**) leads to an outcome (state)
- Agents have preferences (**preferred outcomes**)
- Preferences → outcome **utilities**
- Agents have **degrees of belief** (**probabilities**) for actions

Decision theory = **probability theory** + **utility theory**

Decision Theory

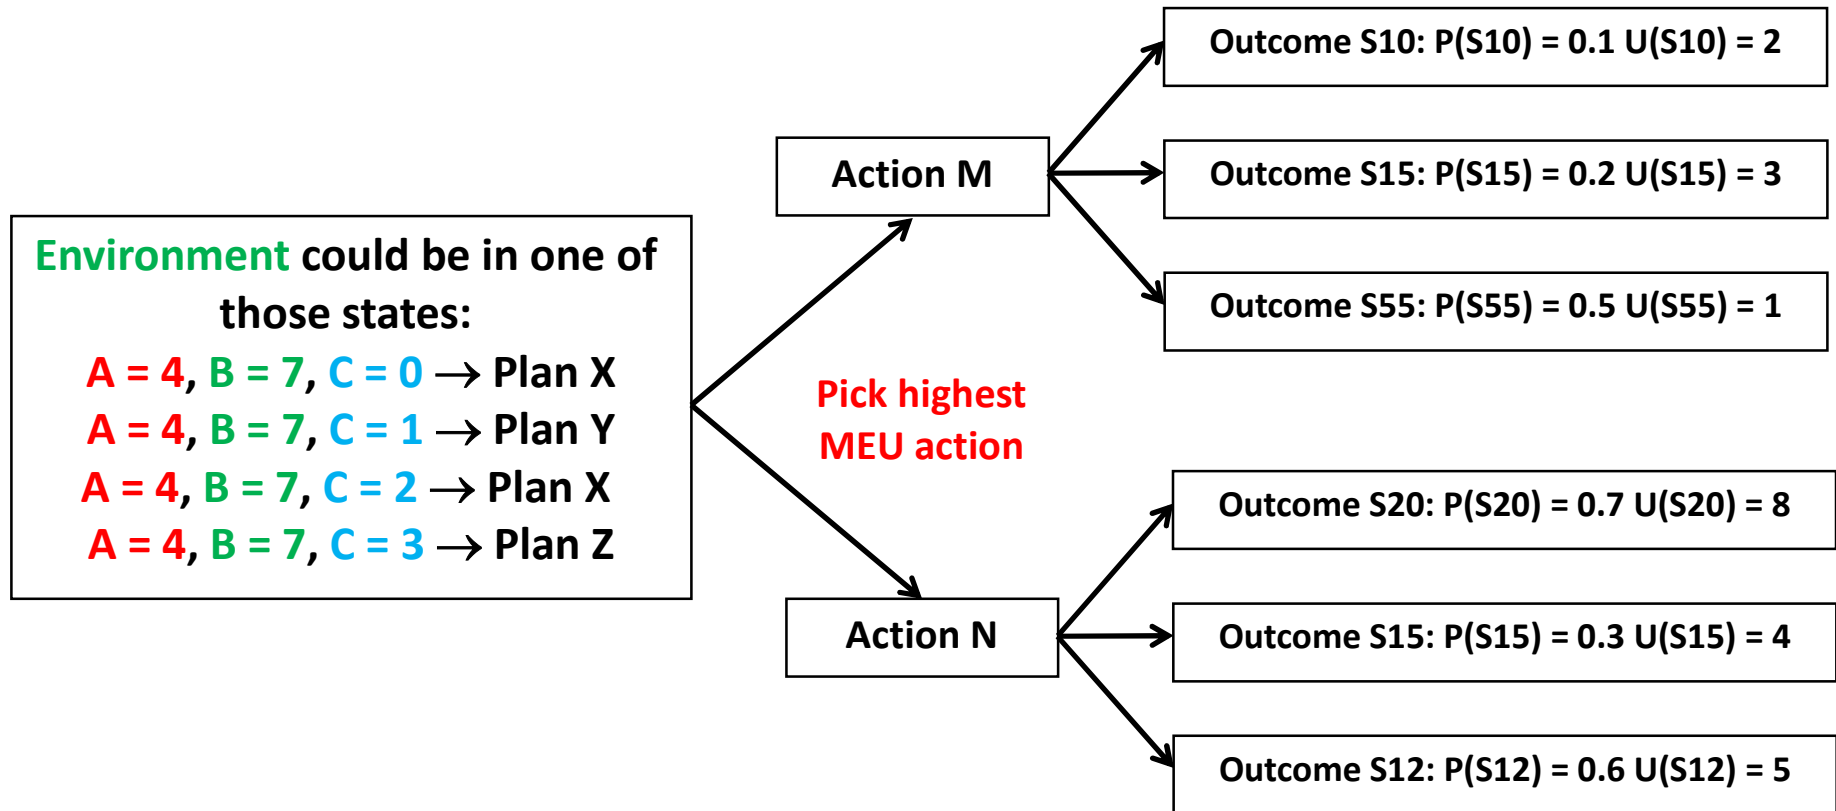
- **Decisions**: every plan (**actions**) leads to an outcome (state)
- Agents have preferences (**preferred outcomes**)
- Preferences → outcome **utilities**
- Agents have **degrees of belief** (**probabilities**) for actions

Decision theory = **probability theory** + **utility theory**

BELIEFS **DESIRES**

Maximum Expected (Average) Utility

$$MEU(M) = P(S10) * U(S10) + P(S15) * U(S15) + P(S55) * U(S55)$$



$$MEU(N) = P(S20) * U(S20) + P(S15) * U(S15) + P(S12) * U(S12)$$

Agents Decisions

Recall that agent **ACTIONS** change the state:

- if we are in state **s**
- action **a** is expected to
- lead to another state **s'** (outcome)

Given uncertainty about the current state **s** and action outcome **s'** we need to define the following:

- probability (belief) of being in state **s**: $P(\mathbf{s})$
- probability (belief) of action **a** leading to outcome **s'**: $P(\mathbf{s}' \mid \mathbf{s}, \mathbf{a})$

Now:

$$P(\mathbf{s}' \mid \mathbf{s}, \mathbf{a}) = P(\text{RESULT}(\mathbf{a}) = \mathbf{s}') = \sum_{\mathbf{s}} P(\mathbf{s}) * P(\mathbf{s}' \mid \mathbf{s}, \mathbf{a})$$

State Utility Function

Agent's **preferences (desires)** are captured by the **Utility function** $U(s)$.

Utility function assigns a value to each state s to express how desirable this state is to the agent.

Expected Action Utility

The **expected utility** of an action **a** given the evidence is the **average utility value** of all **possible outcomes s'** of action **a**, **weighted by their probability (belief) of occurrence**:

$$EU(a) = \sum_{s'} \sum_s P(s) * P(s' | s, a) * U(s') = \sum_{s'} P(Result(a) = s') * U(s')$$

Rational agent should choose an action that **maximizes the expected utility**:

$$\text{chosen action} = \underset{a}{\operatorname{argmax}} EU(a)$$

How Did We Get Here?

Let's start with relationships (and related notation) between agent's preferences:

- agent **prefers** A over B:

$$A \sqsupset B$$

- agent is **indifferent** between A and B:

$$A \sim B$$

- agent prefers A over B or is indifferent between A and B (**weak preference**):

$$A \sqsupseteq B$$

The Concept of Lottery

Let's assume the following:

- an **action** a is a lottery ticket
- the **set of outcomes (resulting states)** is a lottery

A lottery L with possible outcomes S_1, \dots, S_n that occur with probabilities p_1, \dots, p_n is written as:

$$L = [p_1, S_1; p_2, S_2; \dots ; p_n, S_n]$$

Lottery outcome S_i : atomic state or another lottery.

Lottery Constraints: Orderability

Given two lotteries A and B , a rational agent must either prefer one or else rate them as equally preferable:

Exactly one of $(A \sqsubset B)$, $(B \sqsubset A)$, or $(A \sim B)$ holds

Lottery Constraints: Transitivity

Given three lotteries A, B, and C, if an agent prefers A to B AND prefers B to C, then the agent must prefer A to C:

$$(A \sqsubset B) \wedge (B \sqsubset C) \Rightarrow (A \sqsubset C)$$

Lottery Constraints: Continuity

If some lottery B is between A and C in preference, then there is some probability p for which the rational agent will be indifferent between getting B for sure or some other lottery that yields A with probability p and C with probability $1 - p$:

$$(A \sqsubset B \sqsubset C) \Rightarrow \exists p [p, A; 1-p, C] \sim B$$

Lottery Constraints: Substitutability

If an agent is indifferent between two lotteries A and B, then the agent is indifferent between two more complex lotteries that are the same, except that B is substituted for A in one of them:

$$(A \sim B) \Rightarrow [p, A; 1-p, C] \sim [p, B; 1-p, C]$$

Lottery Constraints: Monotonicity


Suppose two lotteries have the same two possible outcomes, A and B. If an agent prefers A to B, then the agent must prefer the lottery that has a higher probability for A:

$$(A \succ B) \Rightarrow (p > q \Leftrightarrow [p, A; 1-p, B] \succ [q, A; 1-q, B])$$

Lottery Constraints: Decomposability

Compound lotteries can be reduced to smaller ones using the laws of probability:

$$[p, A; 1-p, [q, B; 1-q, C]] \sim [p, A; (1-p)q, B; (1-p)(1-q), C]$$



Preferences and Utility Function

An agent whose preferences between lotteries follow the set of axioms (**of utility theory**) below:

- Orderability
- Transitivity
- Continuity
- Substitutability
- Monotonicity
- Decomposability

can be described as possessing a utility function and maximize it.

Preferences and Utility Function

If an agent's preferences obey the axioms of utility theory, then there exist a function U such that:

$$U(A) = U(B) \text{ if and only if } (A \sim B)$$

and

$$U(A) > U(B) \text{ if and only if } (A \succ B)$$


Multiattribute Outcomes

Outcomes can be characterized by more than one attribute. Decisions in such cases are handled by Multiattribute Utility Theory.


Attributes: $X = X_1, \dots, X_n$

Assigned values: $x = \langle x_1, \dots, x_n \rangle$


Strict Dominance: Deterministic




Strict Dominance: Deterministic



Strict Dominance: Uncertain



Strict Dominance: Uncertain



Decision Network (Influence Diagram)

Decision networks (also called influence diagrams) are structures / mechanisms for making rational decisions.

Decision networks are based on Bayesian networks, but include additional nodes that represent **actions** and **utilities**.

Decision Networks


The most basic decision network needs to include:

- information about current state s
- possible actions
- resulting state s' (after applying chosen action a)
- utility of the resulting state $U(s')$


Decision Network Nodes

Decision networks are built using the following nodes:


- chance nodes:



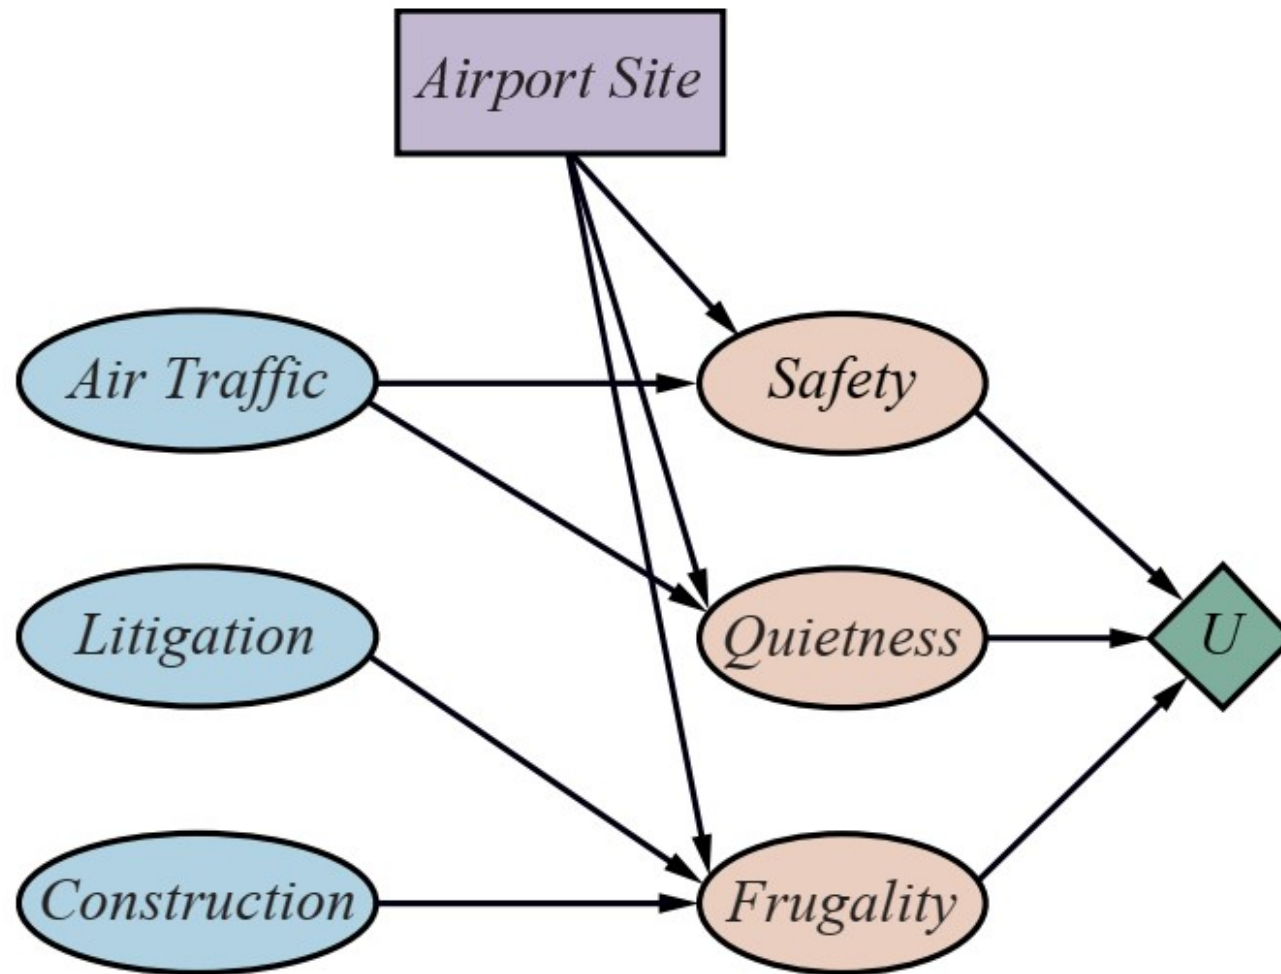
- decision nodes:



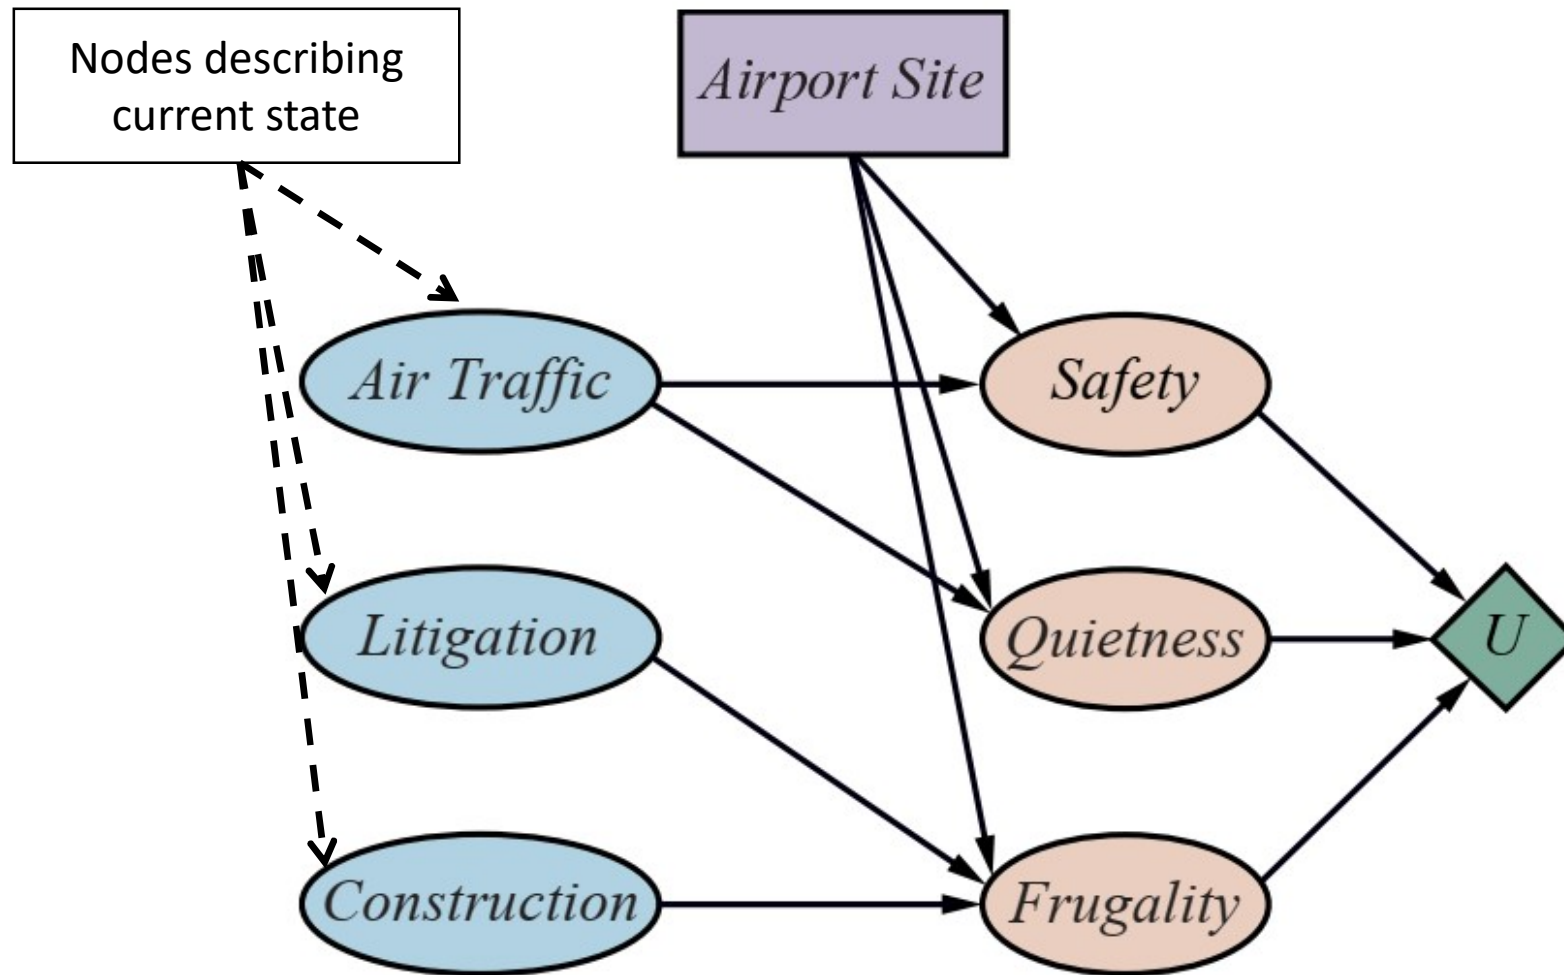
- utility (or value) nodes



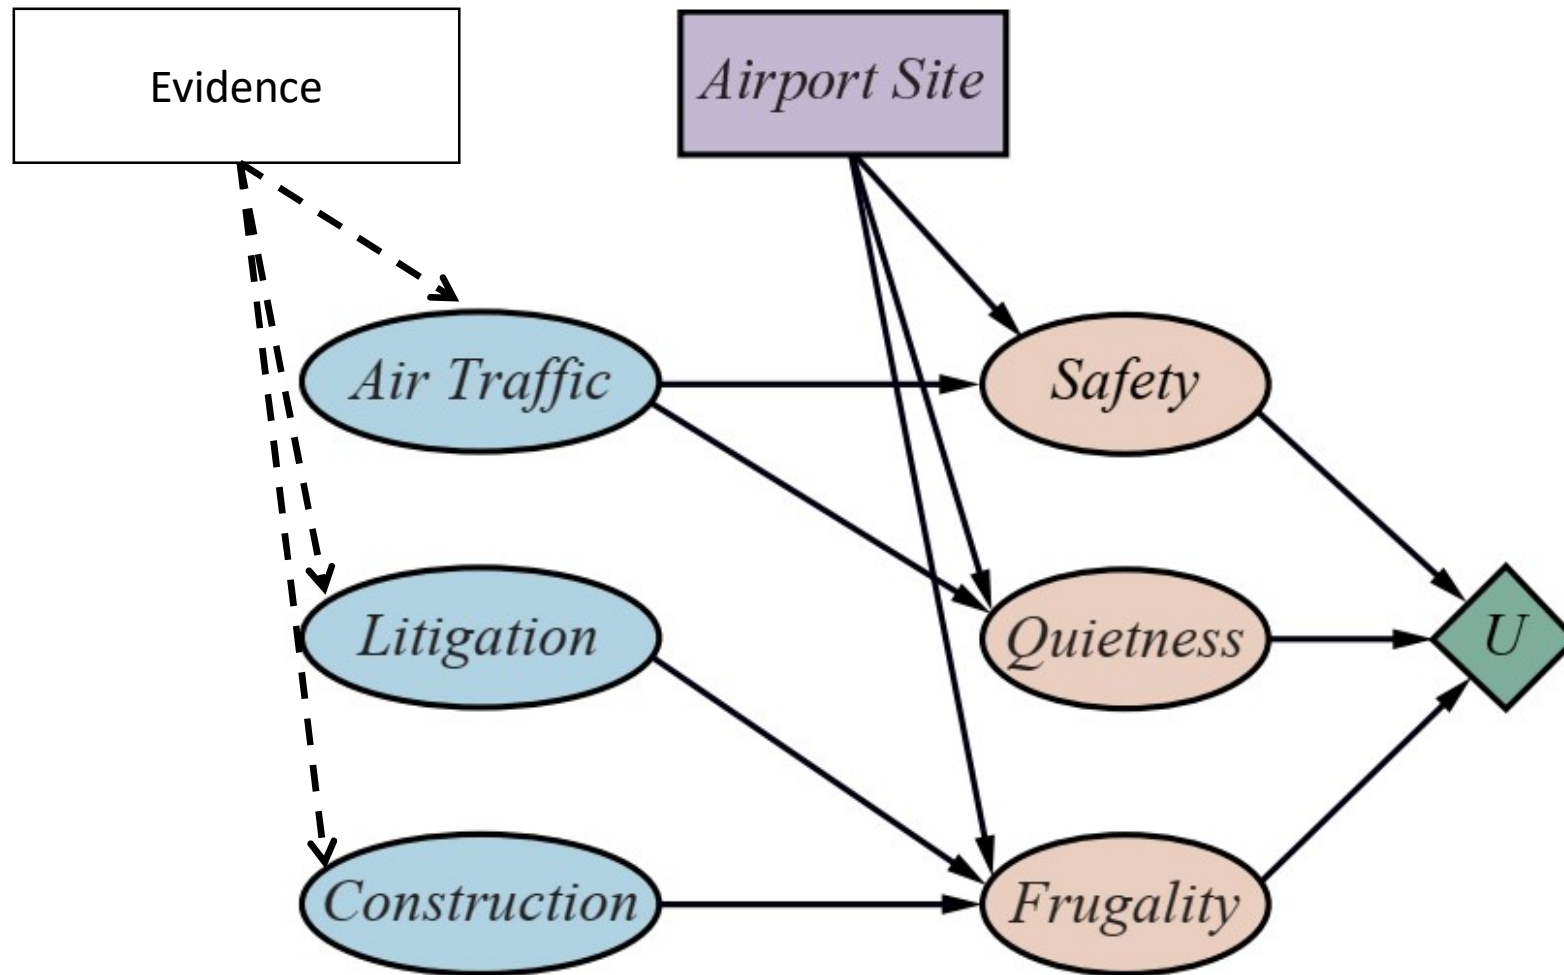
Decision Network: Example



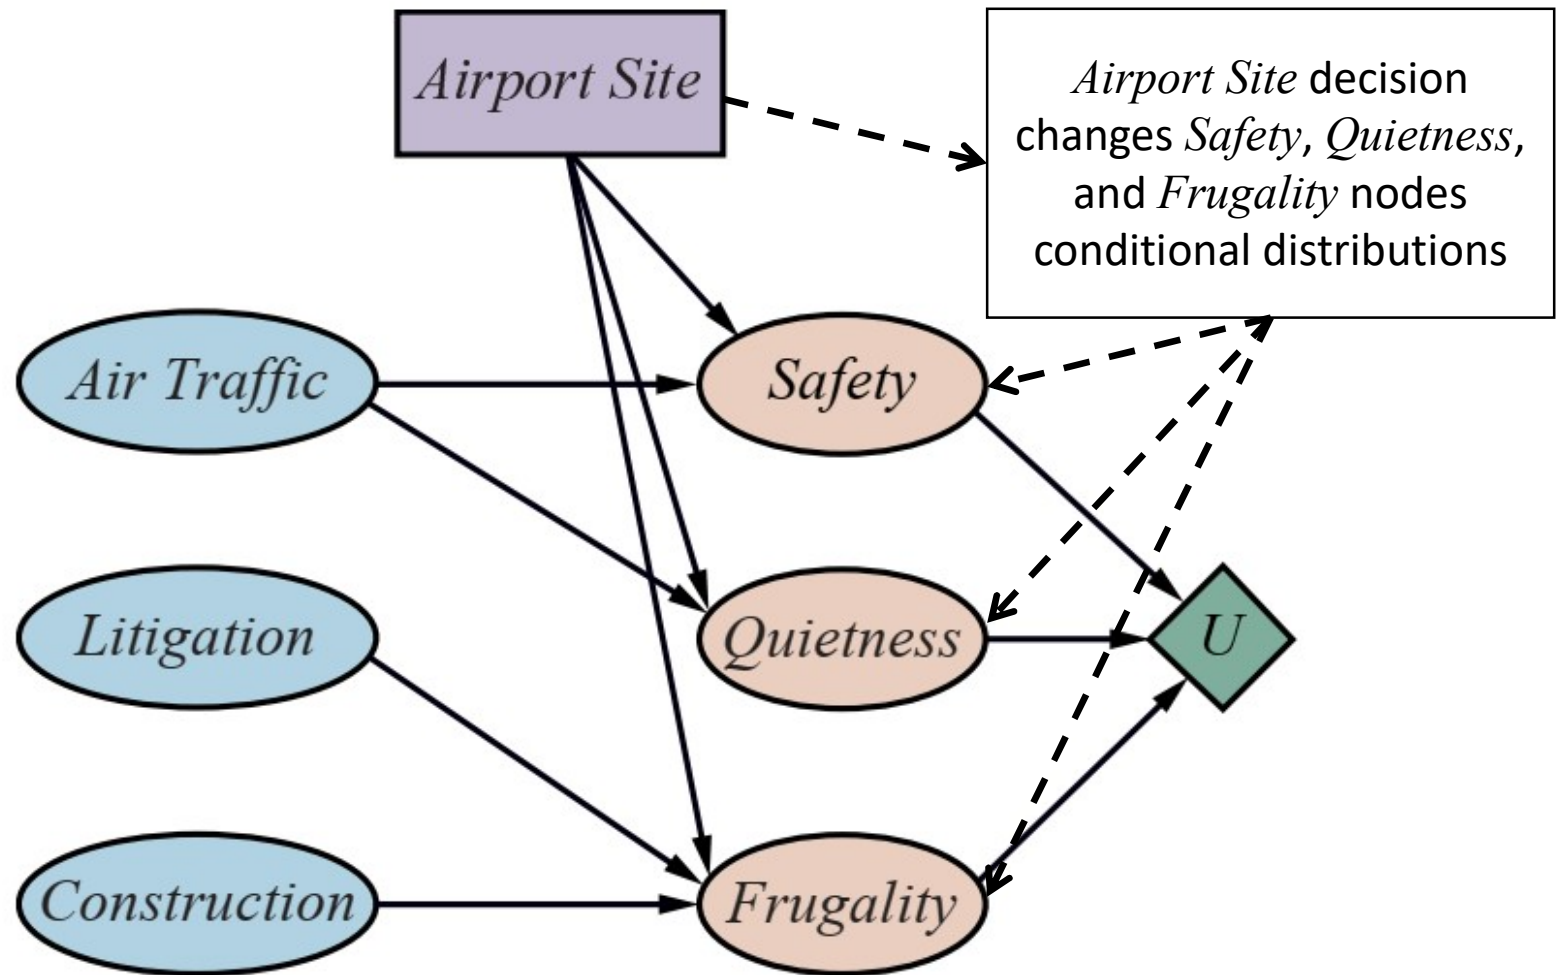
Decision Network: Example



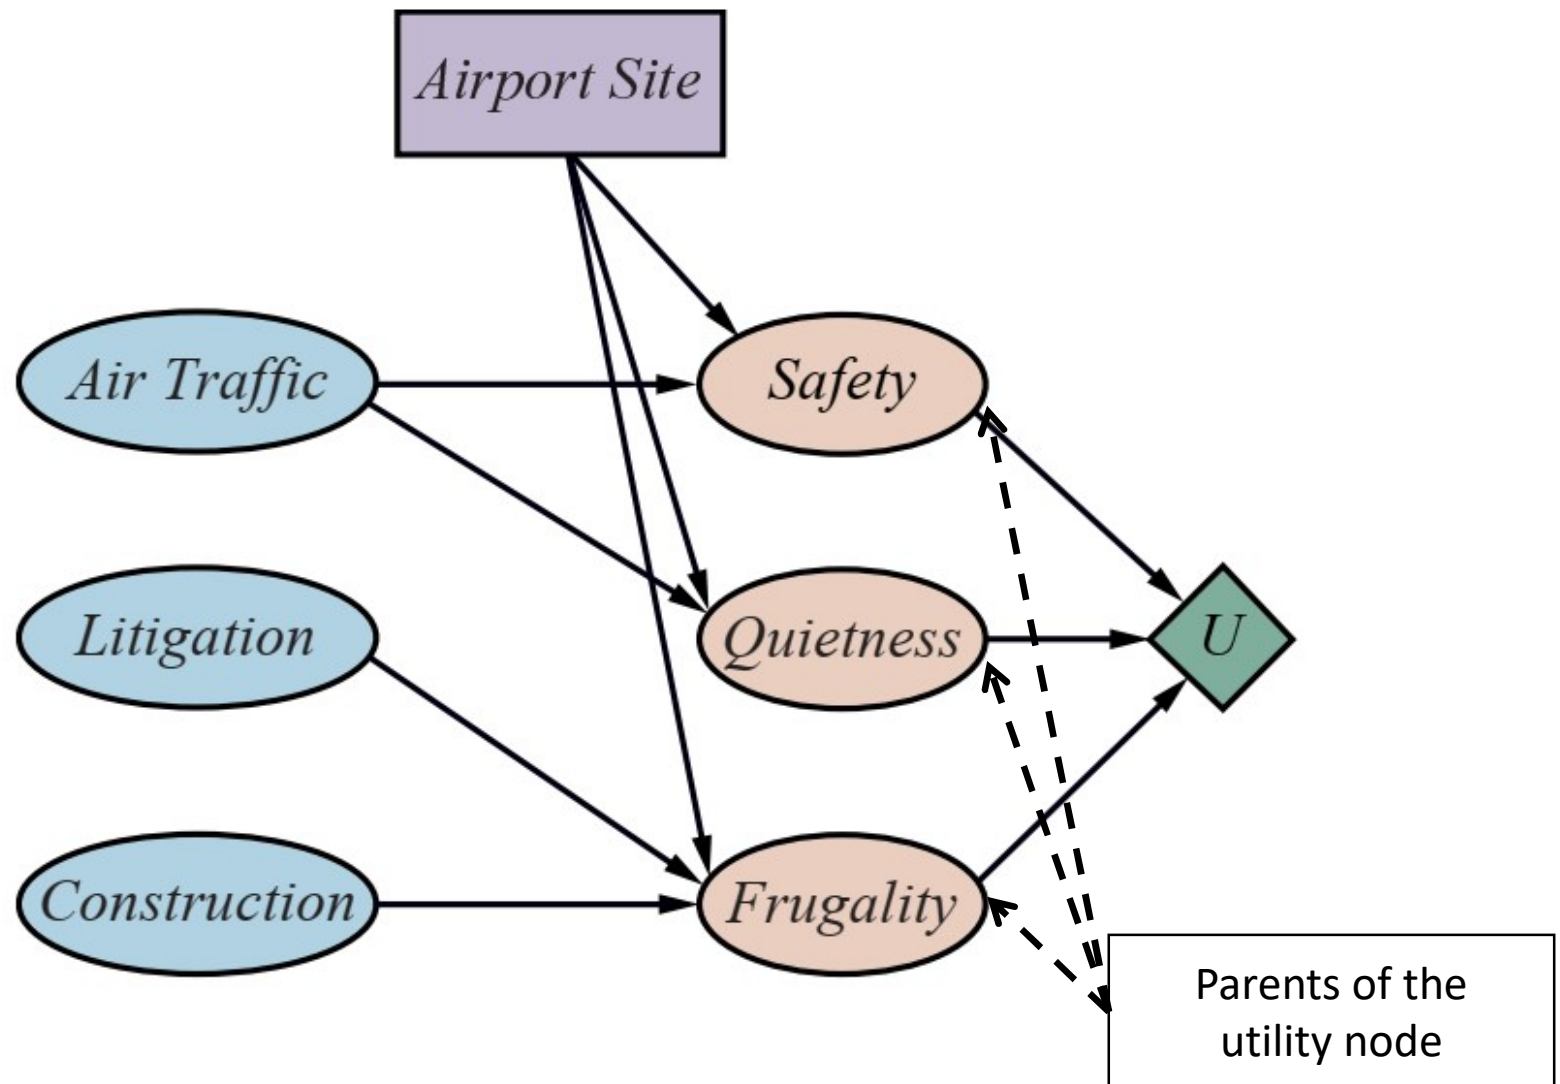
Decision Network: Example



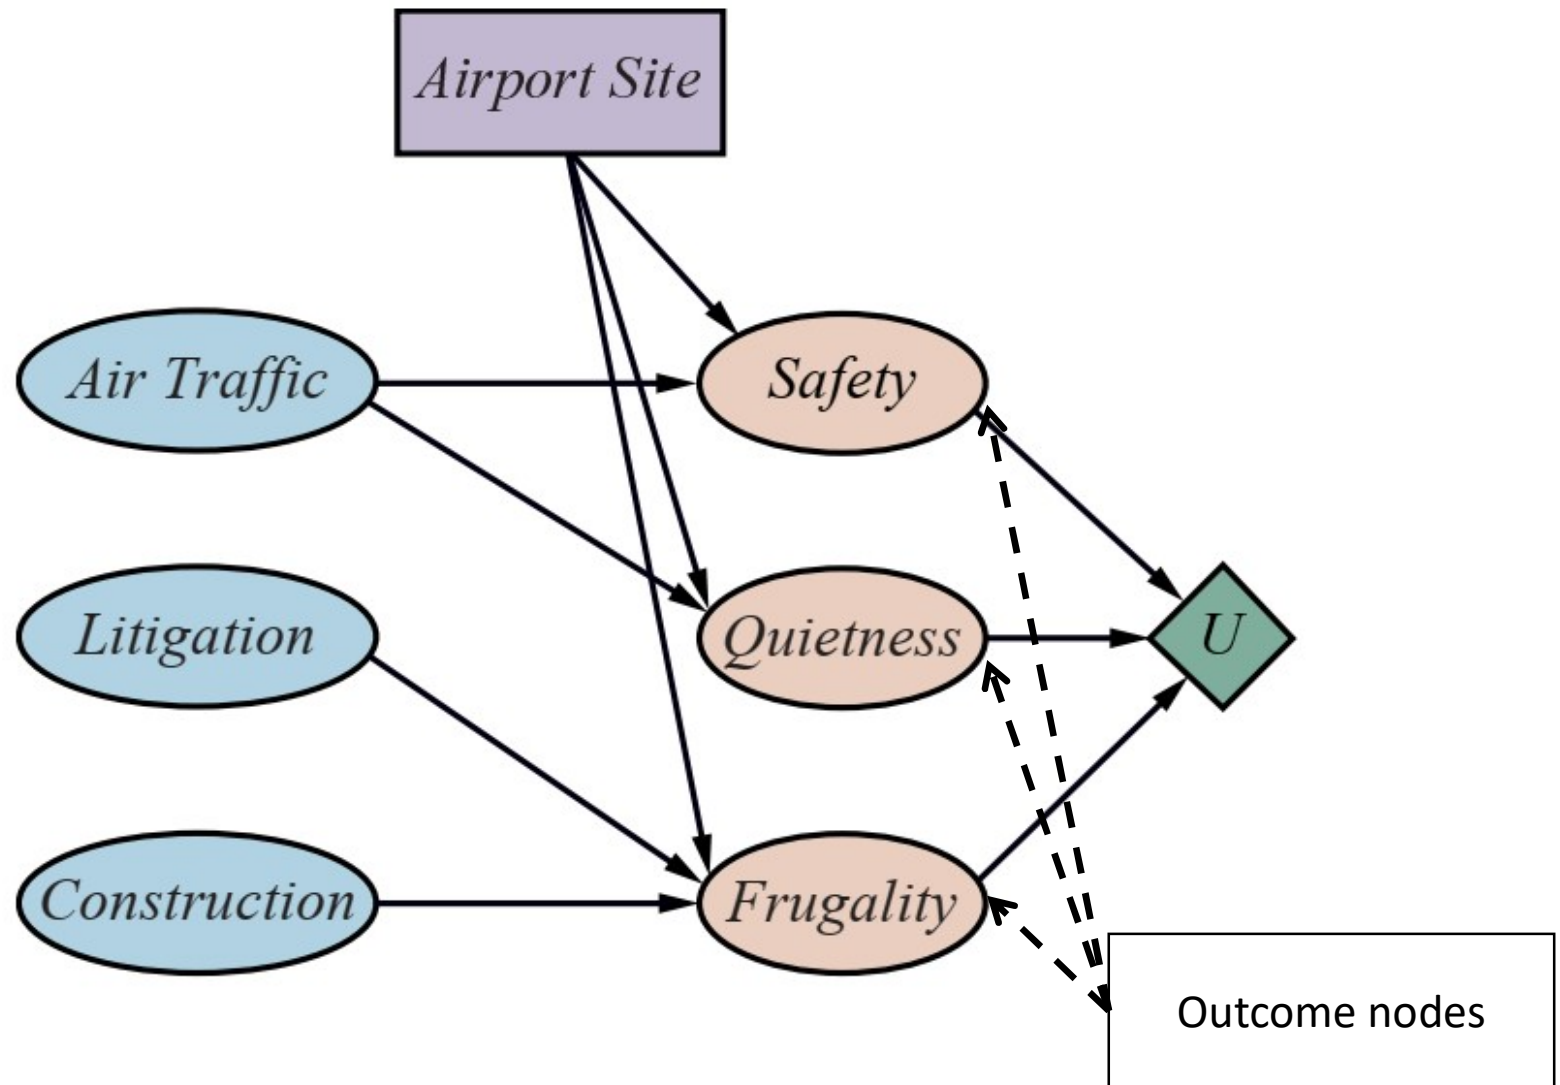
Decision Network: Example



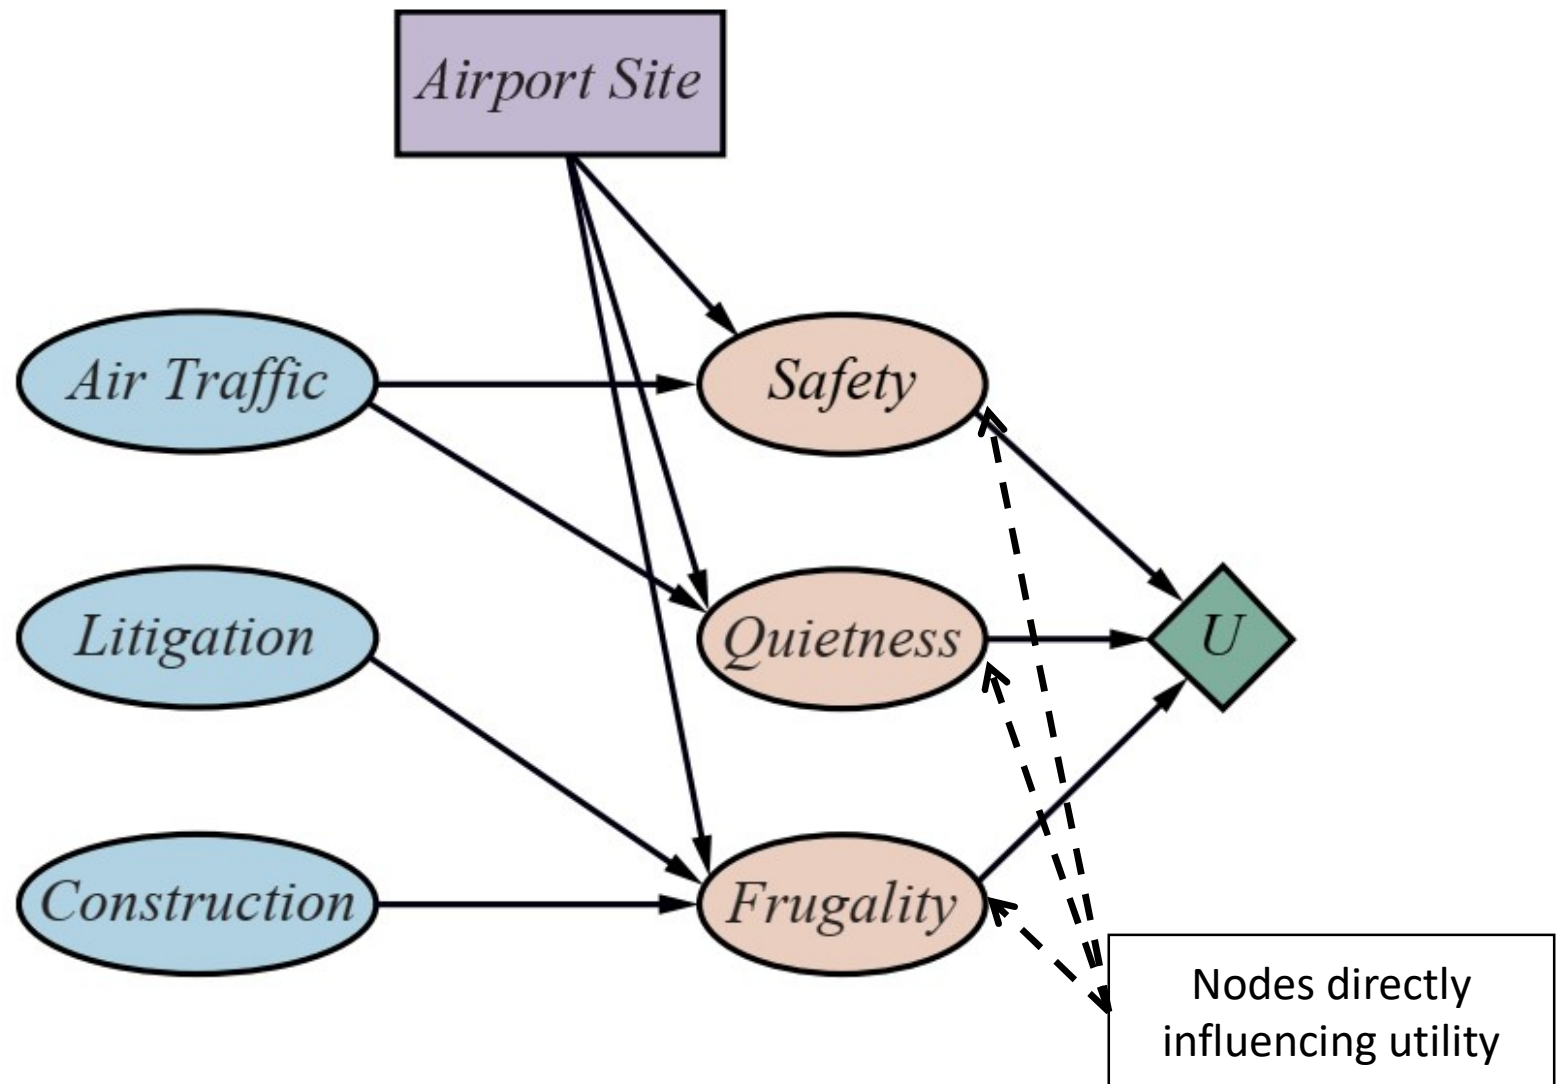
Decision Network: Example



Decision Network: Example



Decision Network: Example

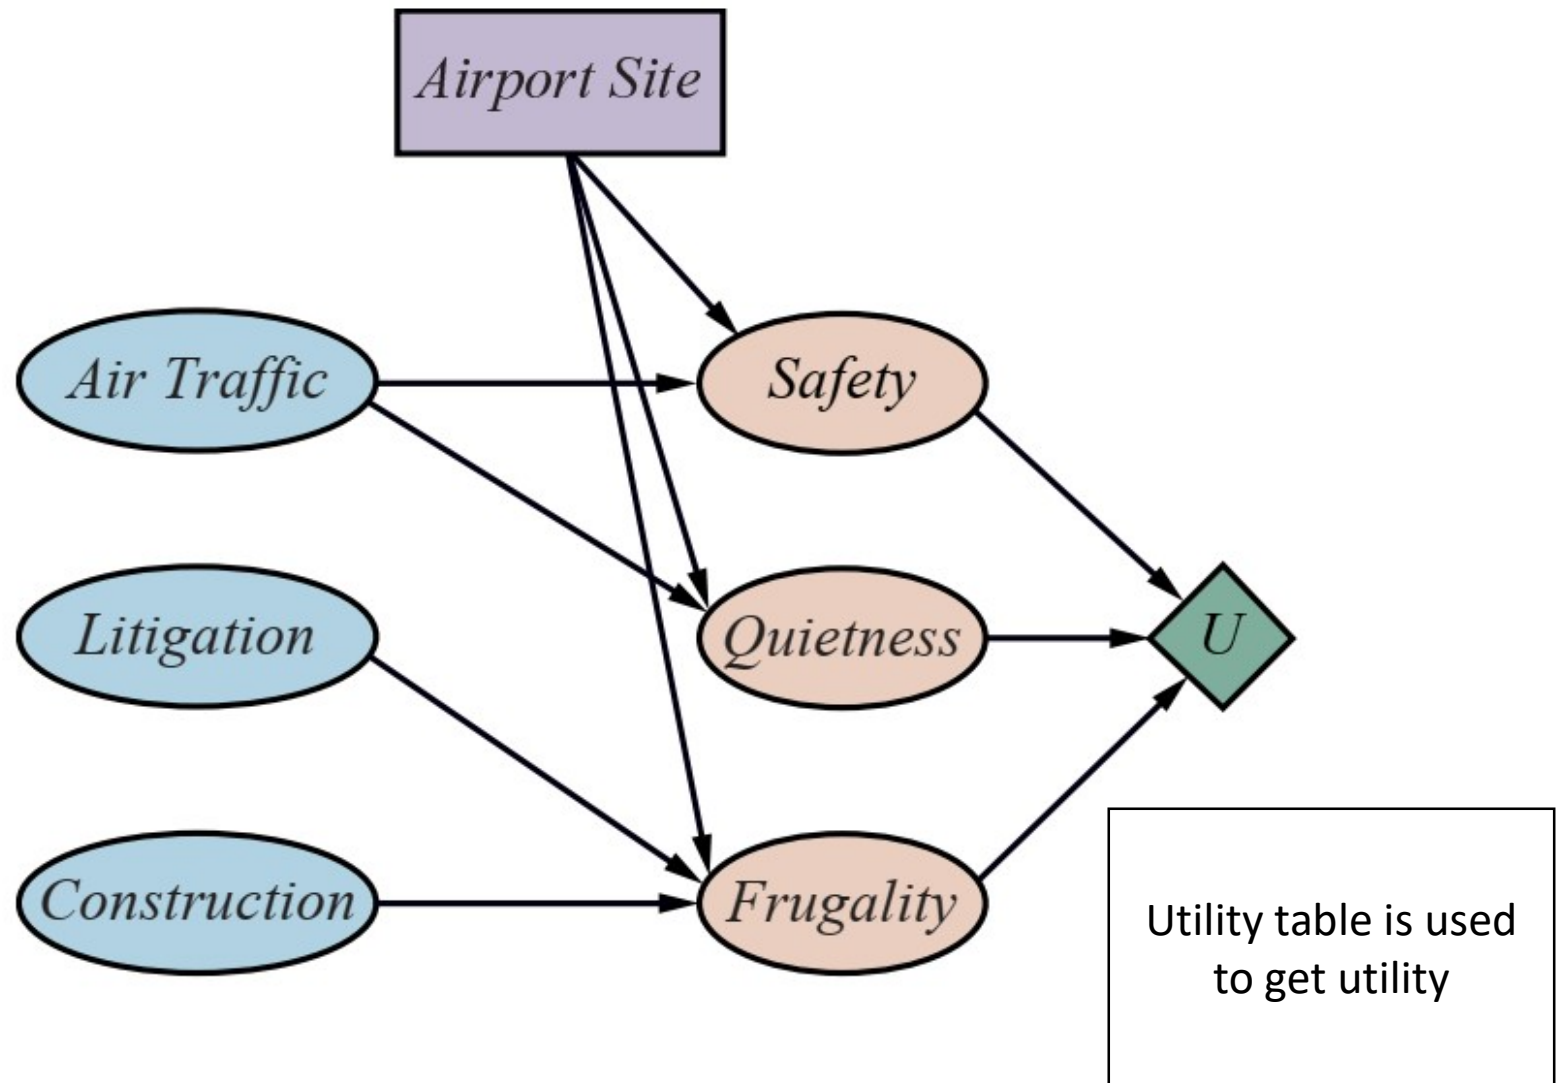


Decision Network: Evaluation

The algorithm for decision network evaluation is as follows:

1. Set the evidence variables for the current state
2. For each possible value a of decision node:
 - a. Set the decision node to that value
 - b. Calculate the posterior probabilities for the parent nodes of the utility node
 - c. Calculate the utility for the action / value a
3. Return the action with highest utility

Decision Network: Example



Decision Network: Simplified Form

