

CS 480

Introduction to Artificial Intelligence

September 15, 2022

Announcements / Reminders

- Please follow the Week 04 To Do List instructions
- Quiz #01 due on Sunday (09/18/22) at 11:00 PM CST
- Written Assignment #01 due on Tuesday (09/20/22) at 11:00 PM CST
- Programming Assignment #1 will be posted within 1.5 - 2 weeks
- **Midterm** Exam (consider fixed):
 - October 13th, 2022 during lecture time

Plan for Today

- **A* Heuristics revisited**
- **Problem Solving: Adversarial Search**

A* Algorithm: Evaluation Function

Calculate / obtain:

$$f(n) = g(\text{State}_n) + h(\text{State}_n)$$

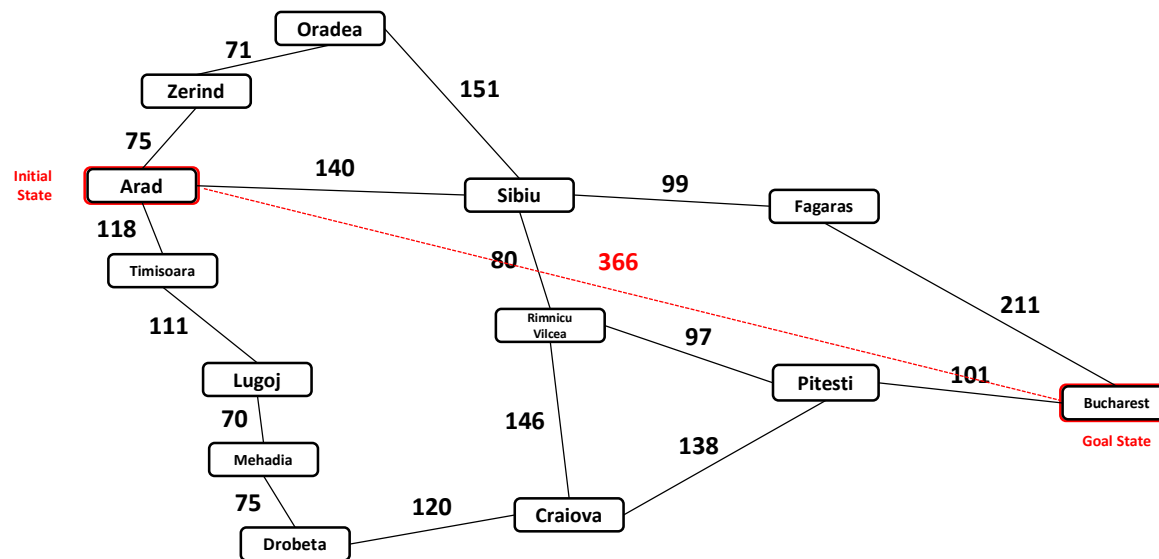
where:

- $g(n)$ - initial node to node n path cost
- $h(n)$ - **estimated cost** of the best path that continues from node n to a goal node

A state n with minimum (maximum) $f(n)$
should be chosen for expansion

What Made A* Work Well?


- Straight-line heuristics is **admissible**: it never overestimates the cost.



- An **admissible heuristics** is guaranteed to give you the optimal solution

What Made A* Work Well?

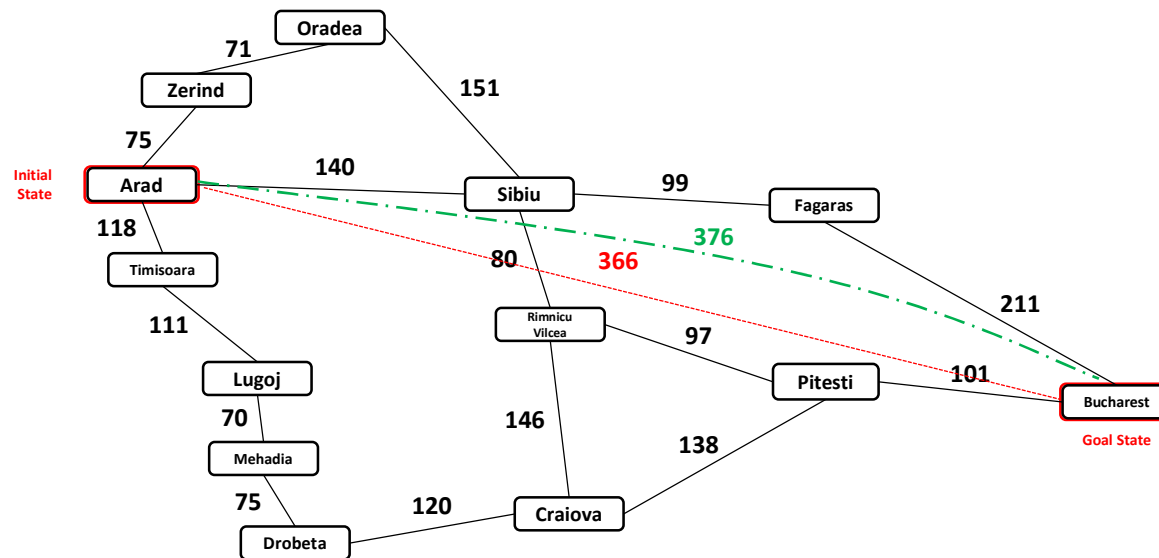
- Straight-line heuristics is **consistent**: its estimate is getting better and better as we get closer to the goal



- Every **consistent** heuristics is **admissible heuristics**, but not the other way around


Dominating Heuristics

- We can have more than one available heuristics. For example $h_1(n)$ and $h_2(n)$.



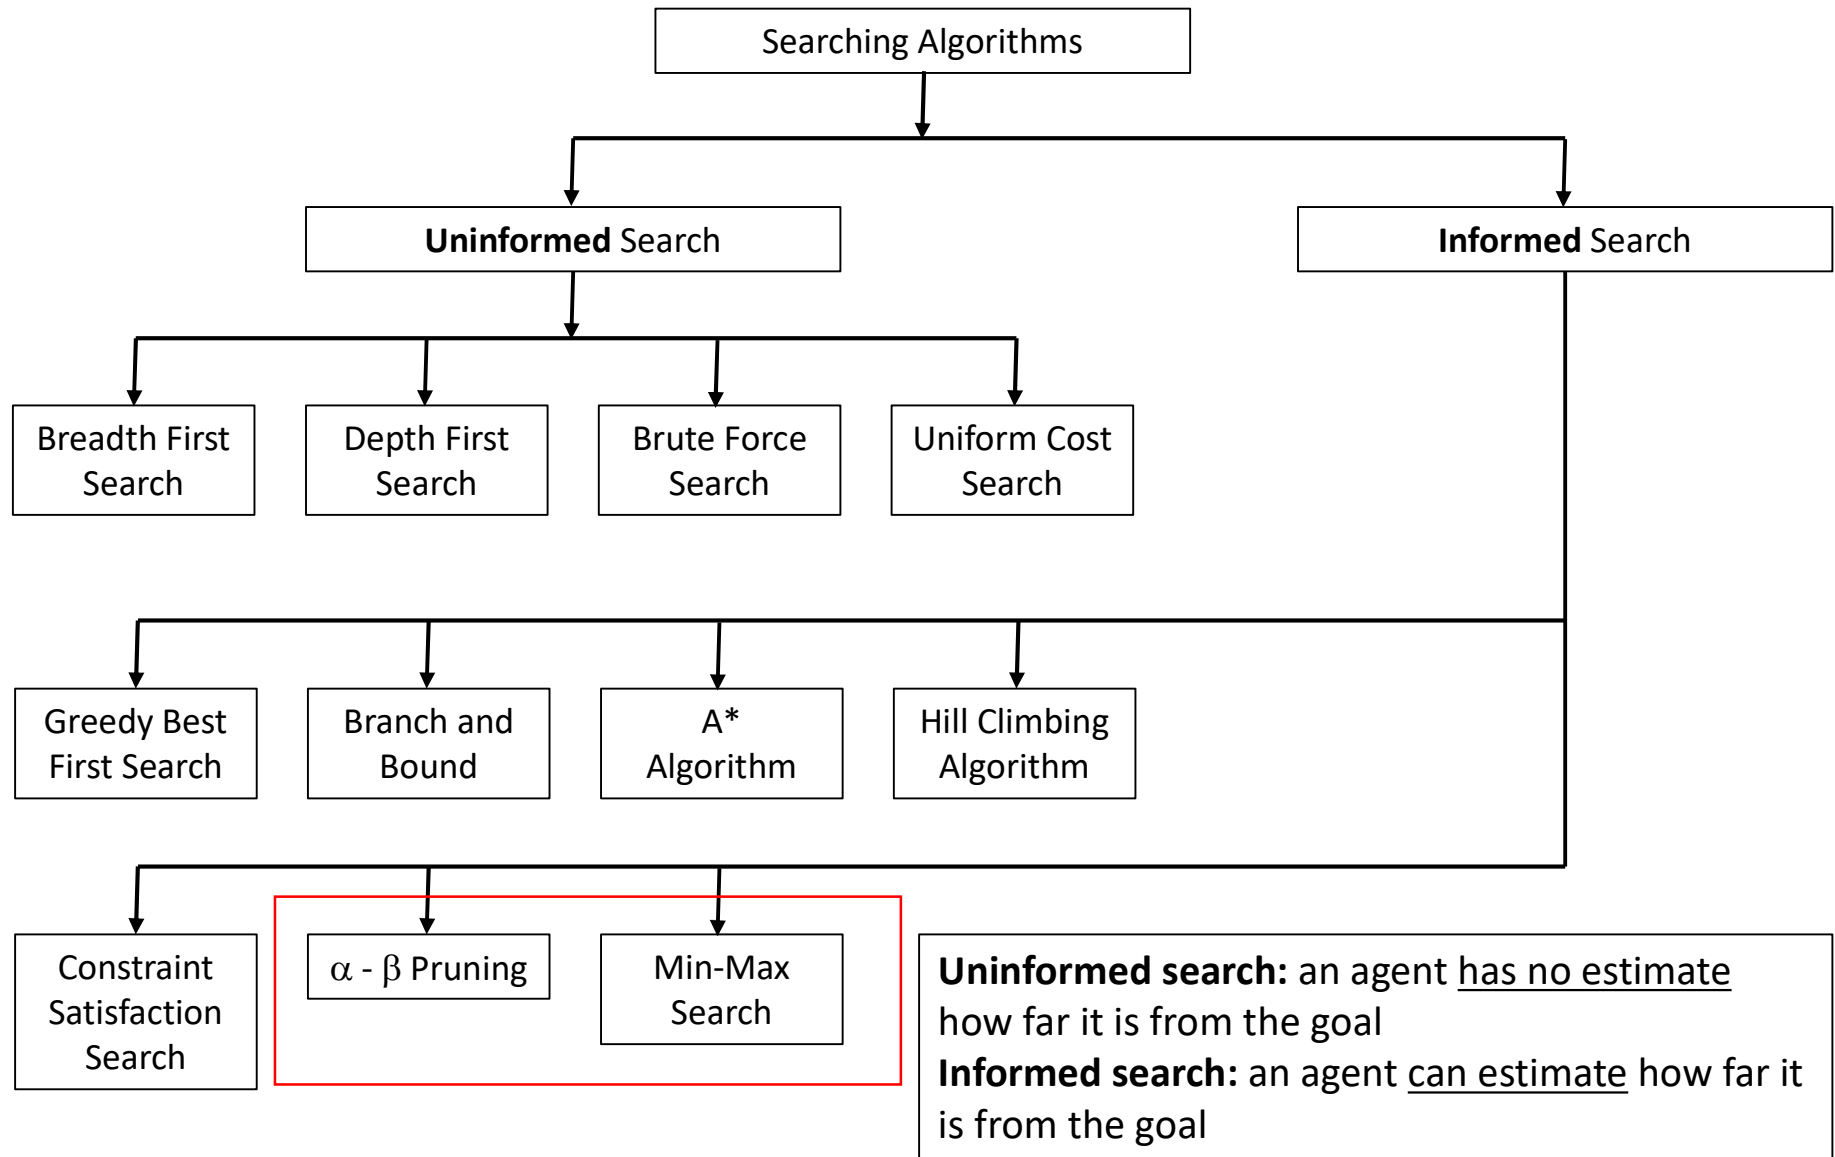
- Heuristics $h_2(n)$ estimate is closer to actual cost than $h_1(n)$. $h_2(n)$ dominates $h_1(n)$. Use $h_2(n)$.

Informed Search: Application Example

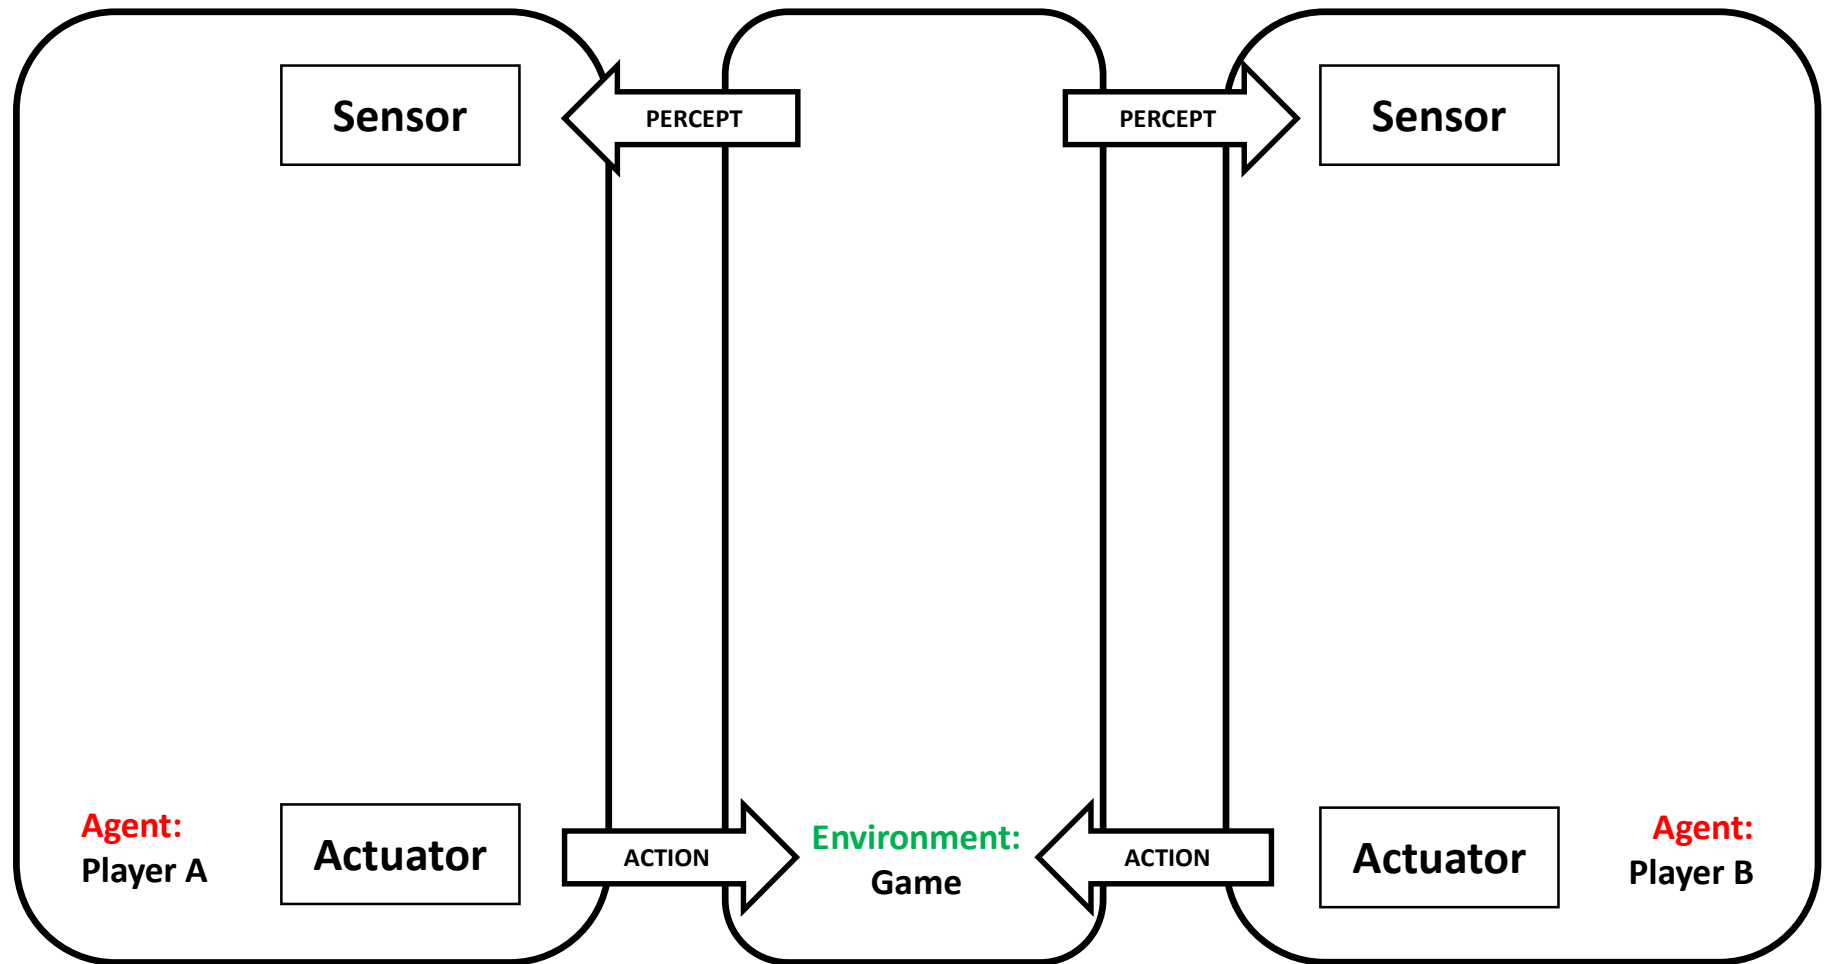


Source: <https://www.youtube.com/watch?v=Pxbv2gEhnMk>

Selected Searching Algorithms



Two-player Games



Perfect Information Zero Sum Games

- Perfect information = fully observable
- Multiagent: number of players is 2 or more
- Multiagent: agents are competitive
- Zero-sum: “winner takes all”
- Examples:
 - Tic Tac Toe
 - Chess

Two Player Games: Env Assumptions

Works with a “Simple Environment”:

- Fully observable
- ~~Single agent~~ Multitagent (competitive!)
- Deterministic
- Static
- Episodic / sequential
- Discrete
- Known to the agent

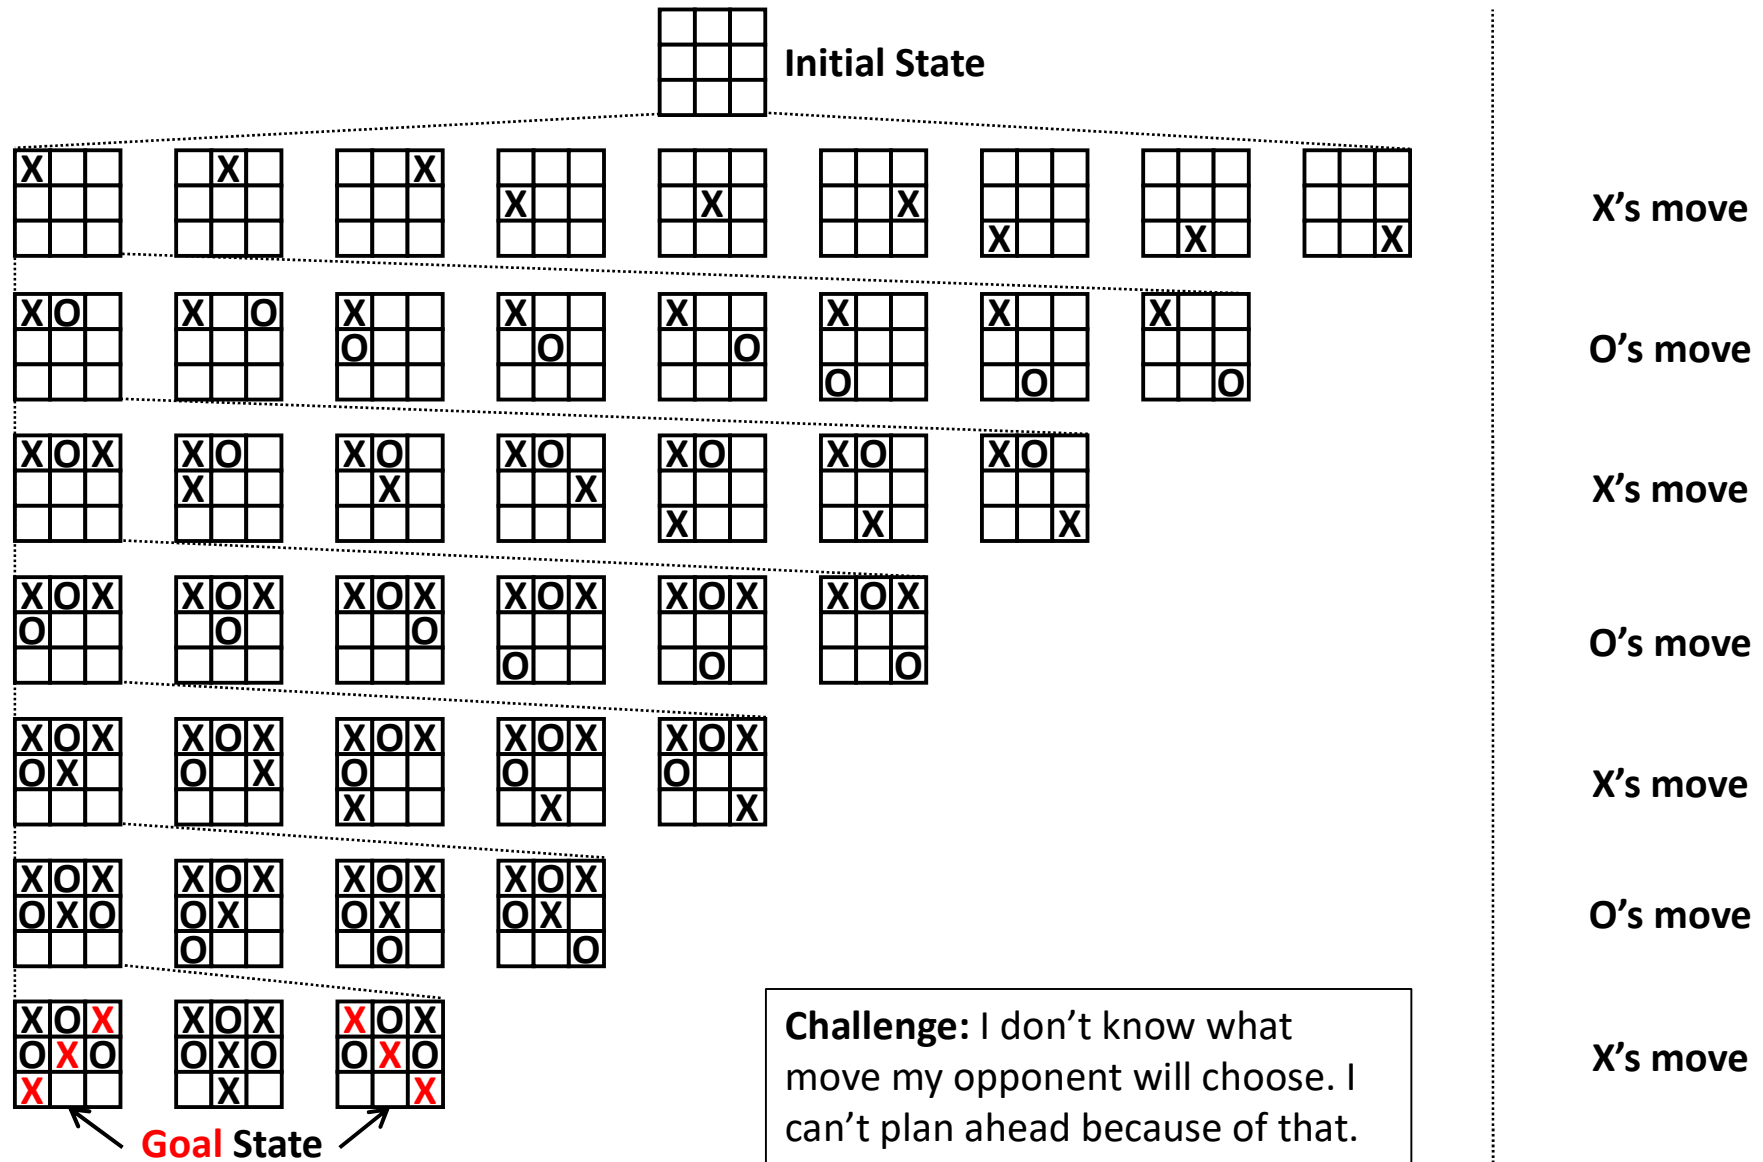
Defining Zero Sum Game Problem

- Define a set of possible states: **State Space**
- Specify how will you track **Whose Move / Turn** it is
- Specify **Initial State**
- Specify **Goal State(s)** (there can be multiple)
- Define a FINITE set of possible **Actions** (legal moves) for EACH state in the State Space
- Come up with a **Transition Model** which describes what each action does
- Come up with a **Terminal Test** that verifies if the game is over
- Specify the **Utility (Payoff / Objective) Function**: a function that defines the final numerical value to player p when the game ends in **terminal state** s


MinMax Algorithm: the Idea

I don't know what move my opponent will choose, but I am going to **ASSUME** that it is going to be the **best / optimal** option


Tic Tac Toe: Zero Sum Game (2 Players)




Tic Tac Toe: Zero Sum Game (2 Players)




State P:
X wins



State R:
Tie




State S:
O wins



Terminal / Leaf States

Tic Tac Toe: Zero Sum Game (2 Players)



State P:
X wins

UTILITY(P) =
1.0

State R:
Tie

UTILITY(R) =
0.0

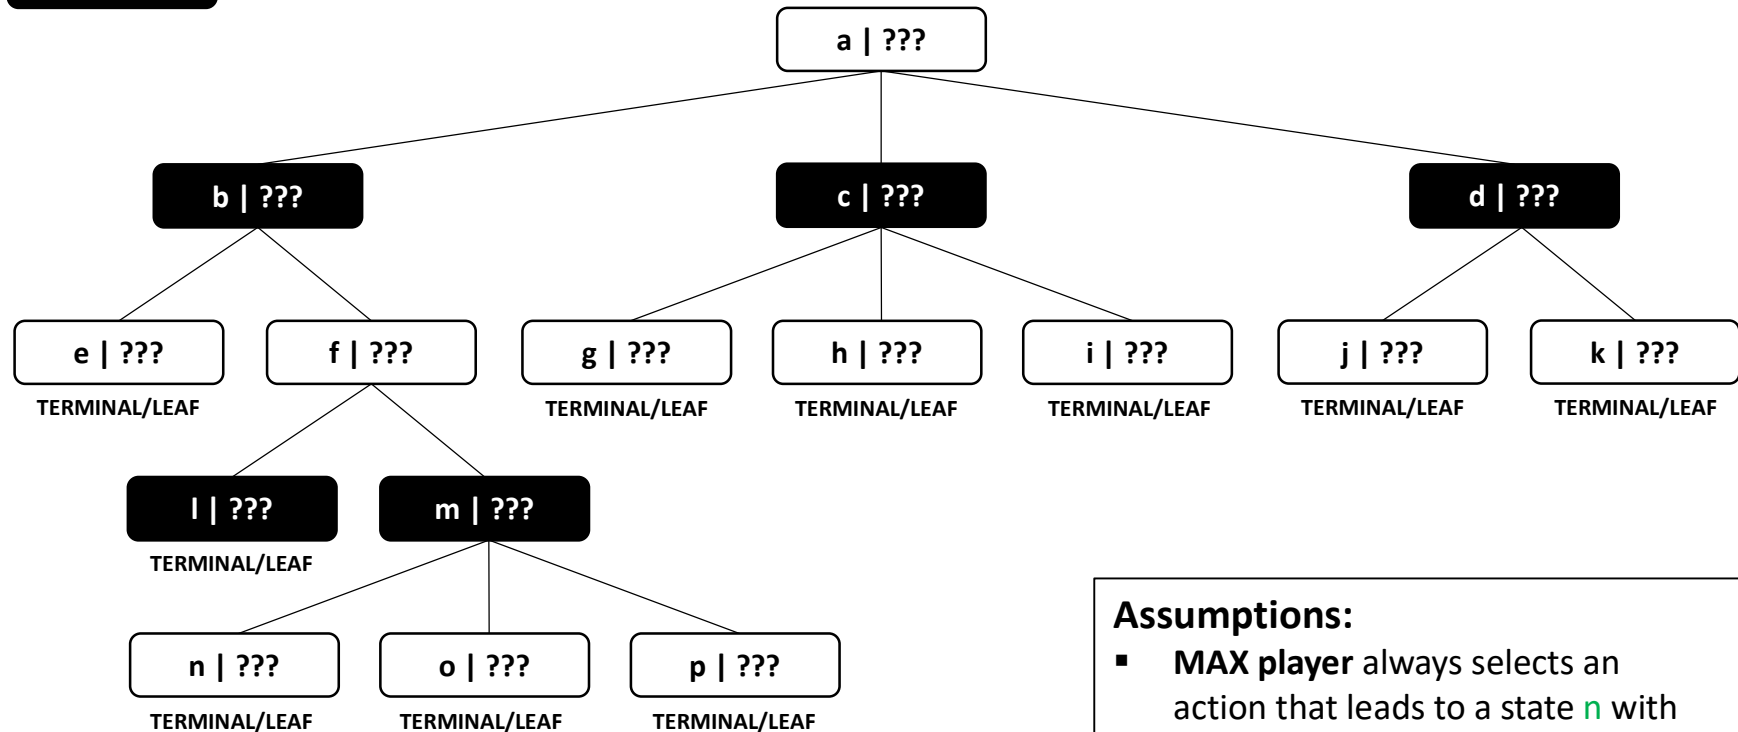
State S:
O wins

UTILITY(S) =
-1.0

Example MinMax Search Tree

$n \mid \text{MINMAX}(n)$ MAX player state / move / turn

$n \mid \text{MINMAX}(n)$ MIN player state / move / turn



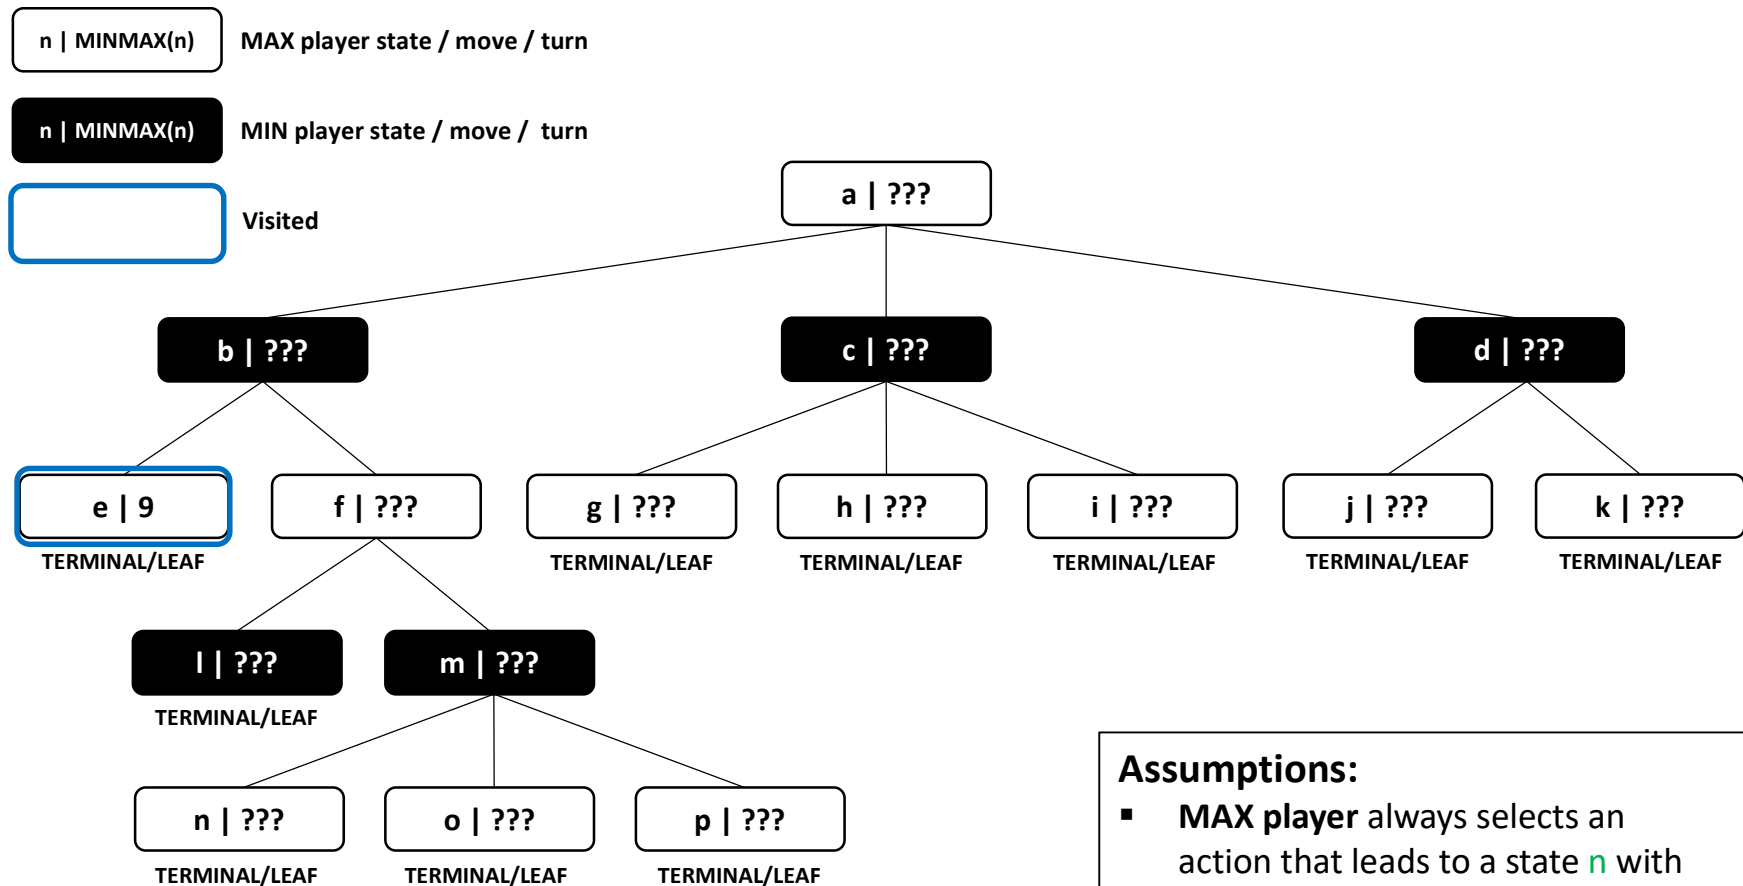
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



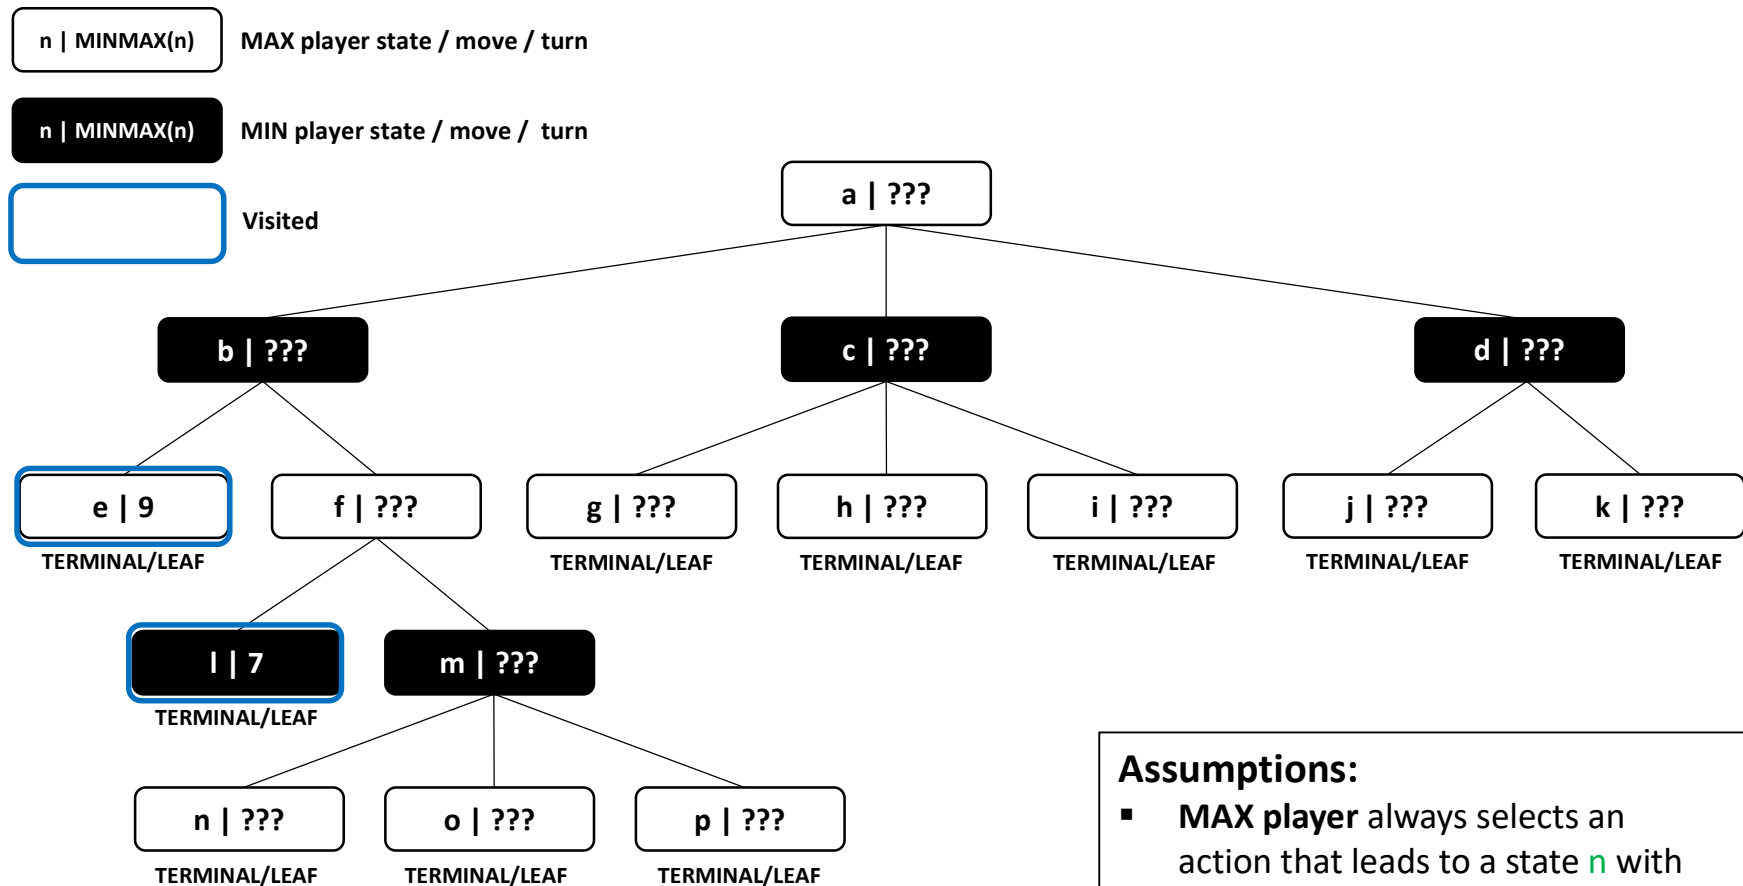
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- MAX player always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- MIN player always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- BOTH players always play optimally

Example MinMax Search Tree



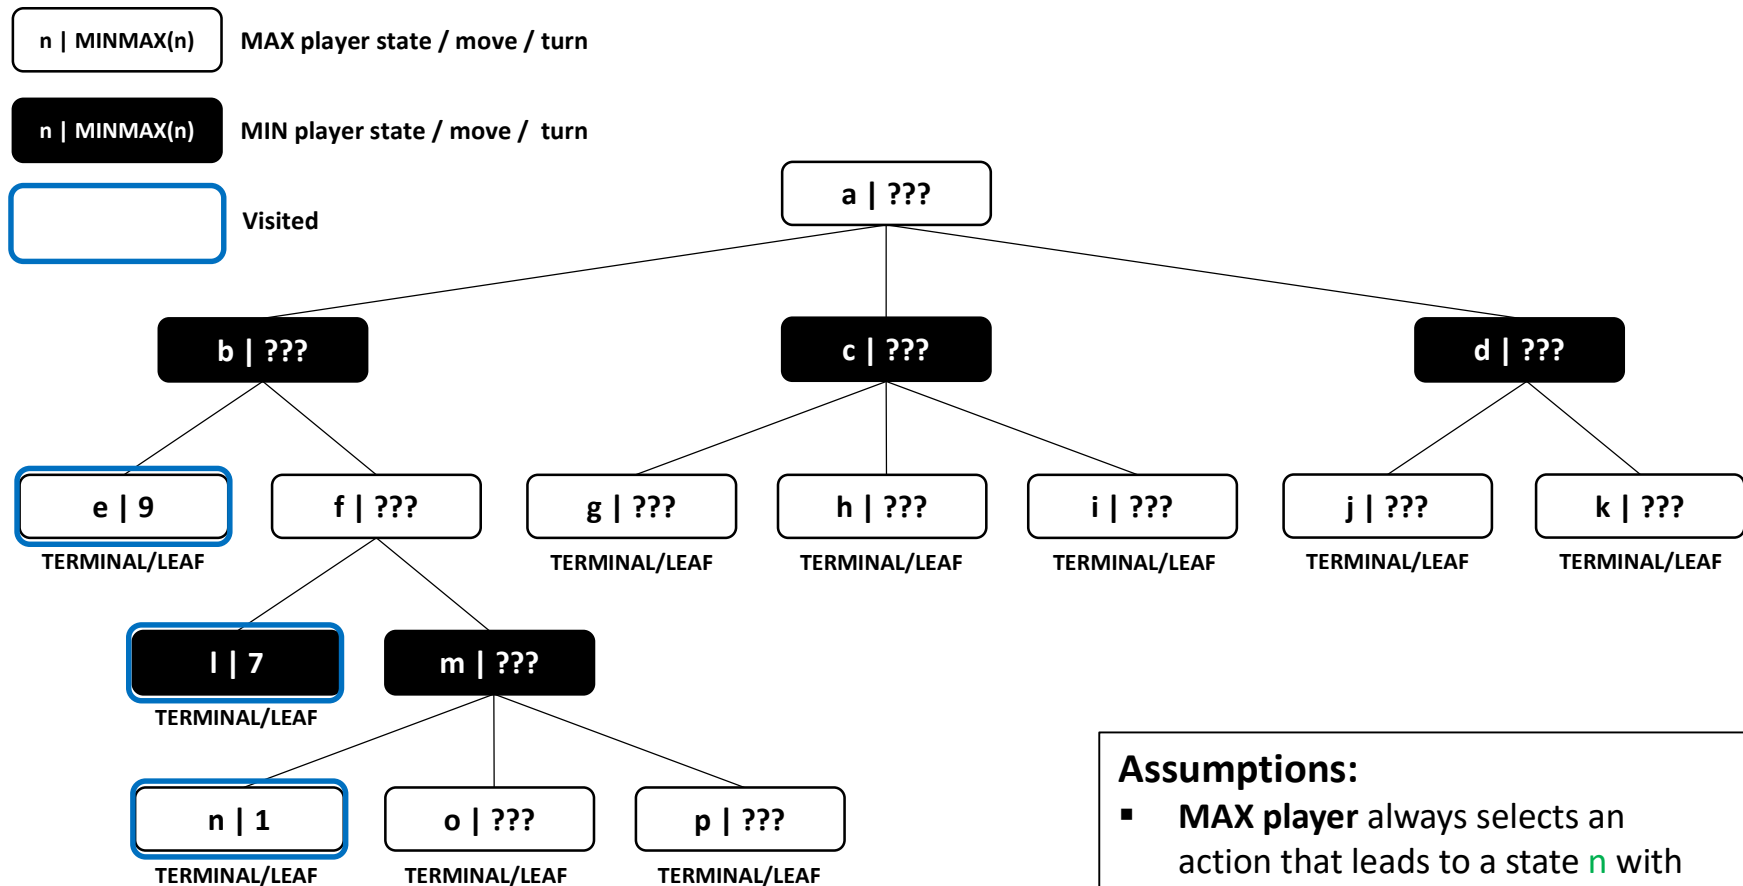
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



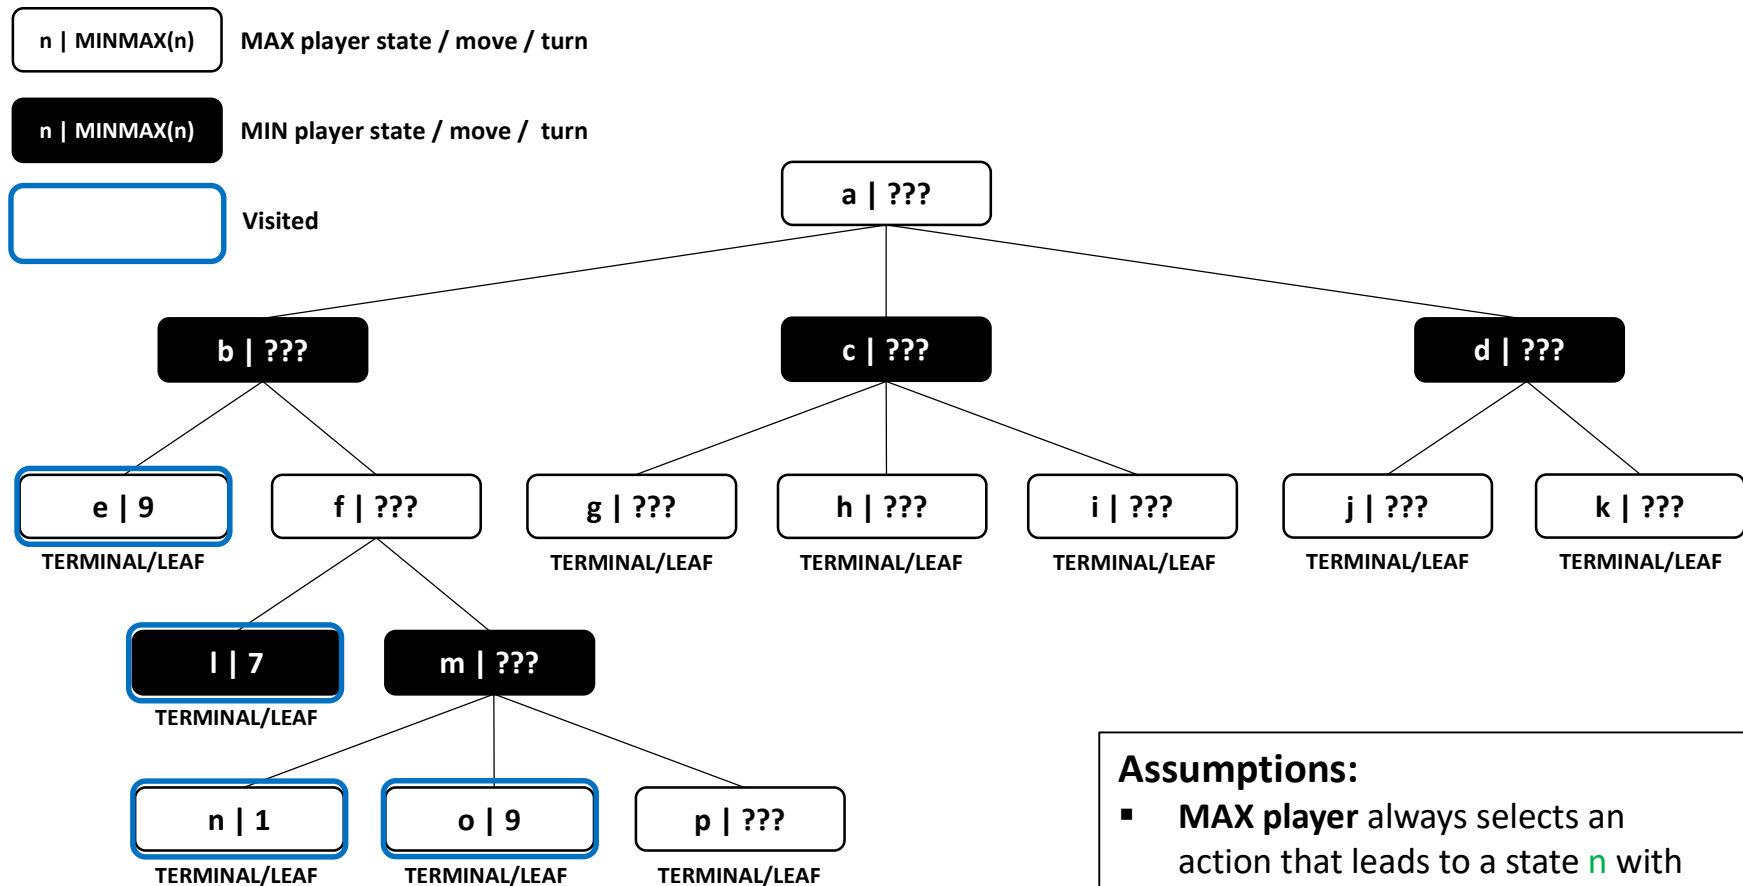
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



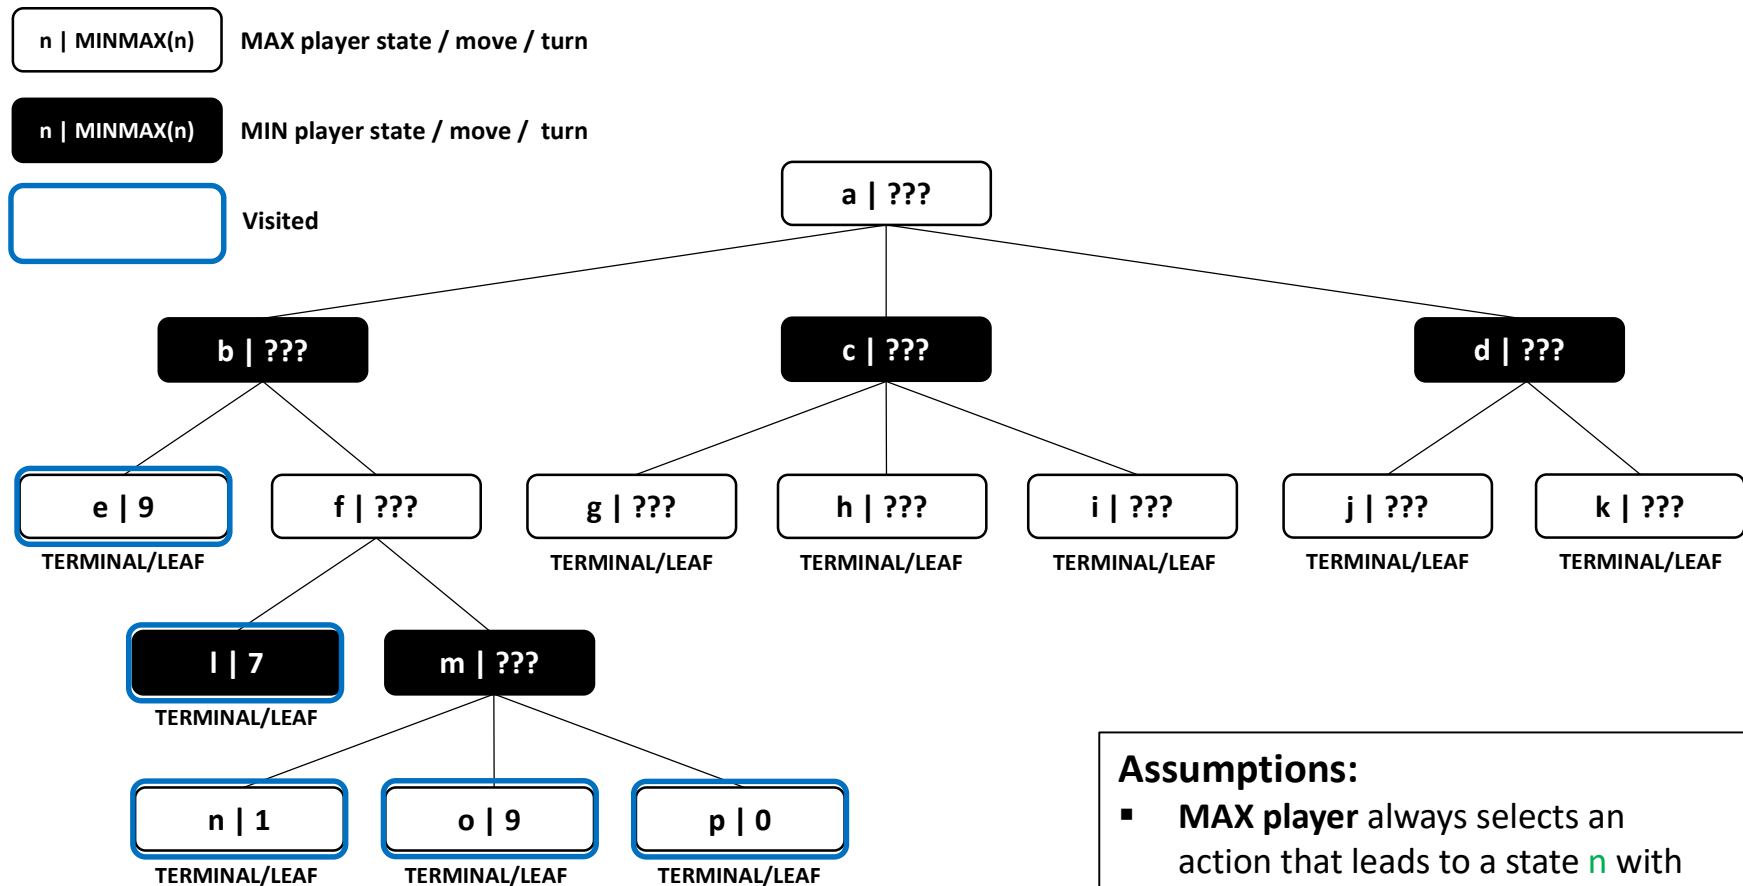
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



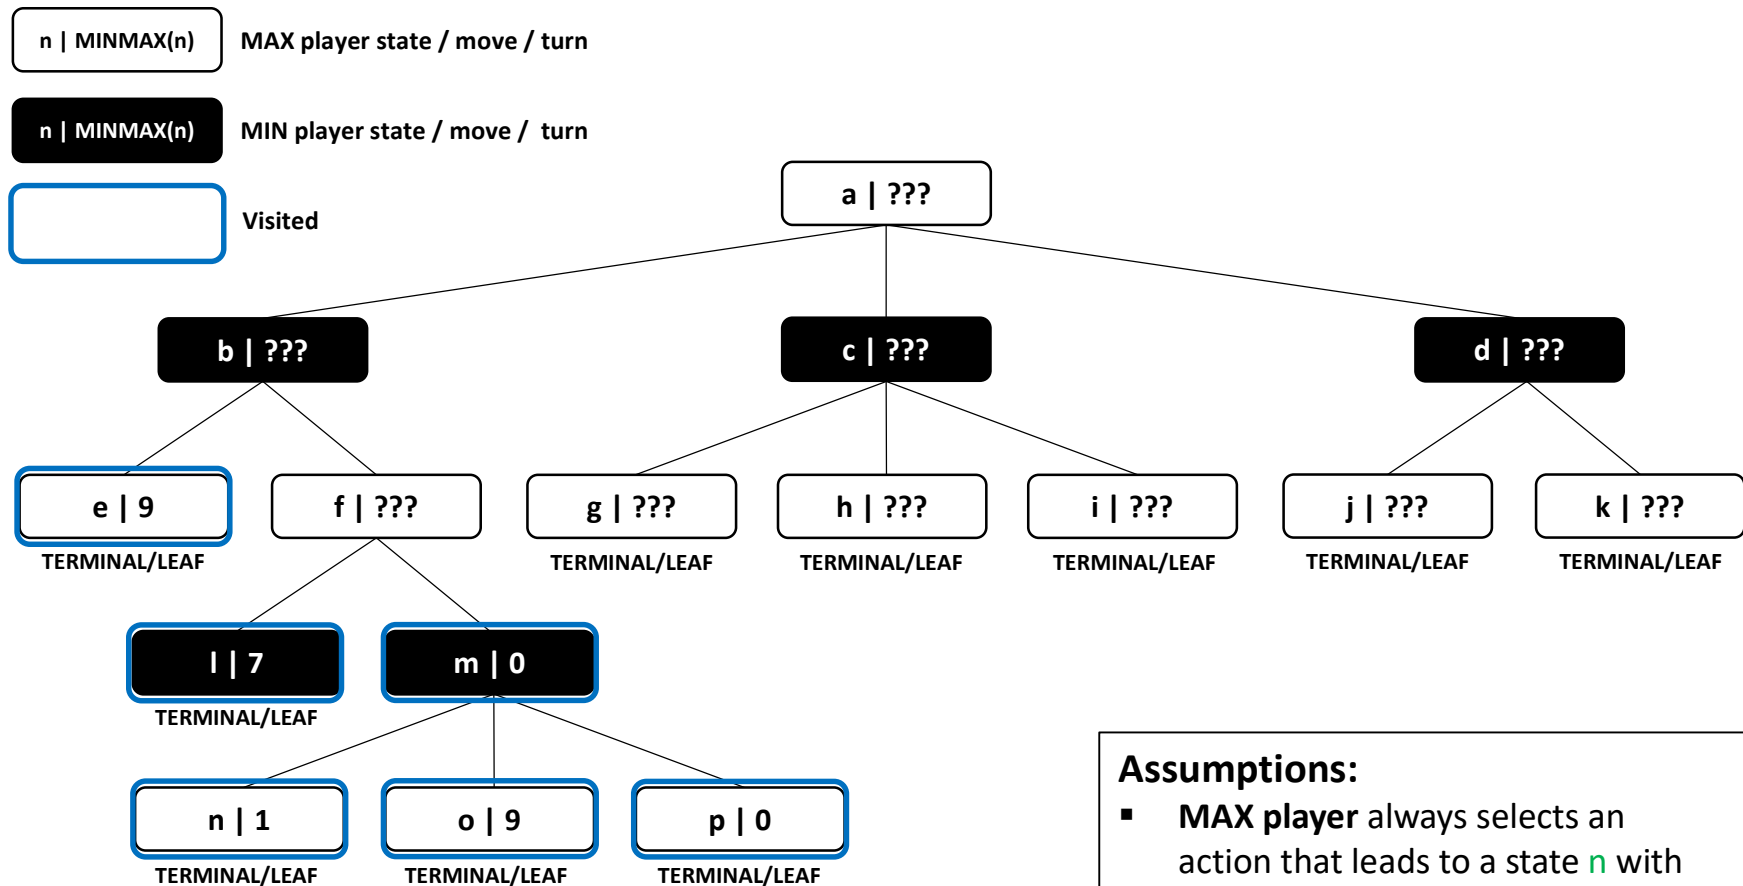
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** MINIMAX(n) value
- **MIN player** always selects an action that leads to a state n with **minimum** MINIMAX(n) value
- **BOTH players** always play optimally

Example MinMax Search Tree



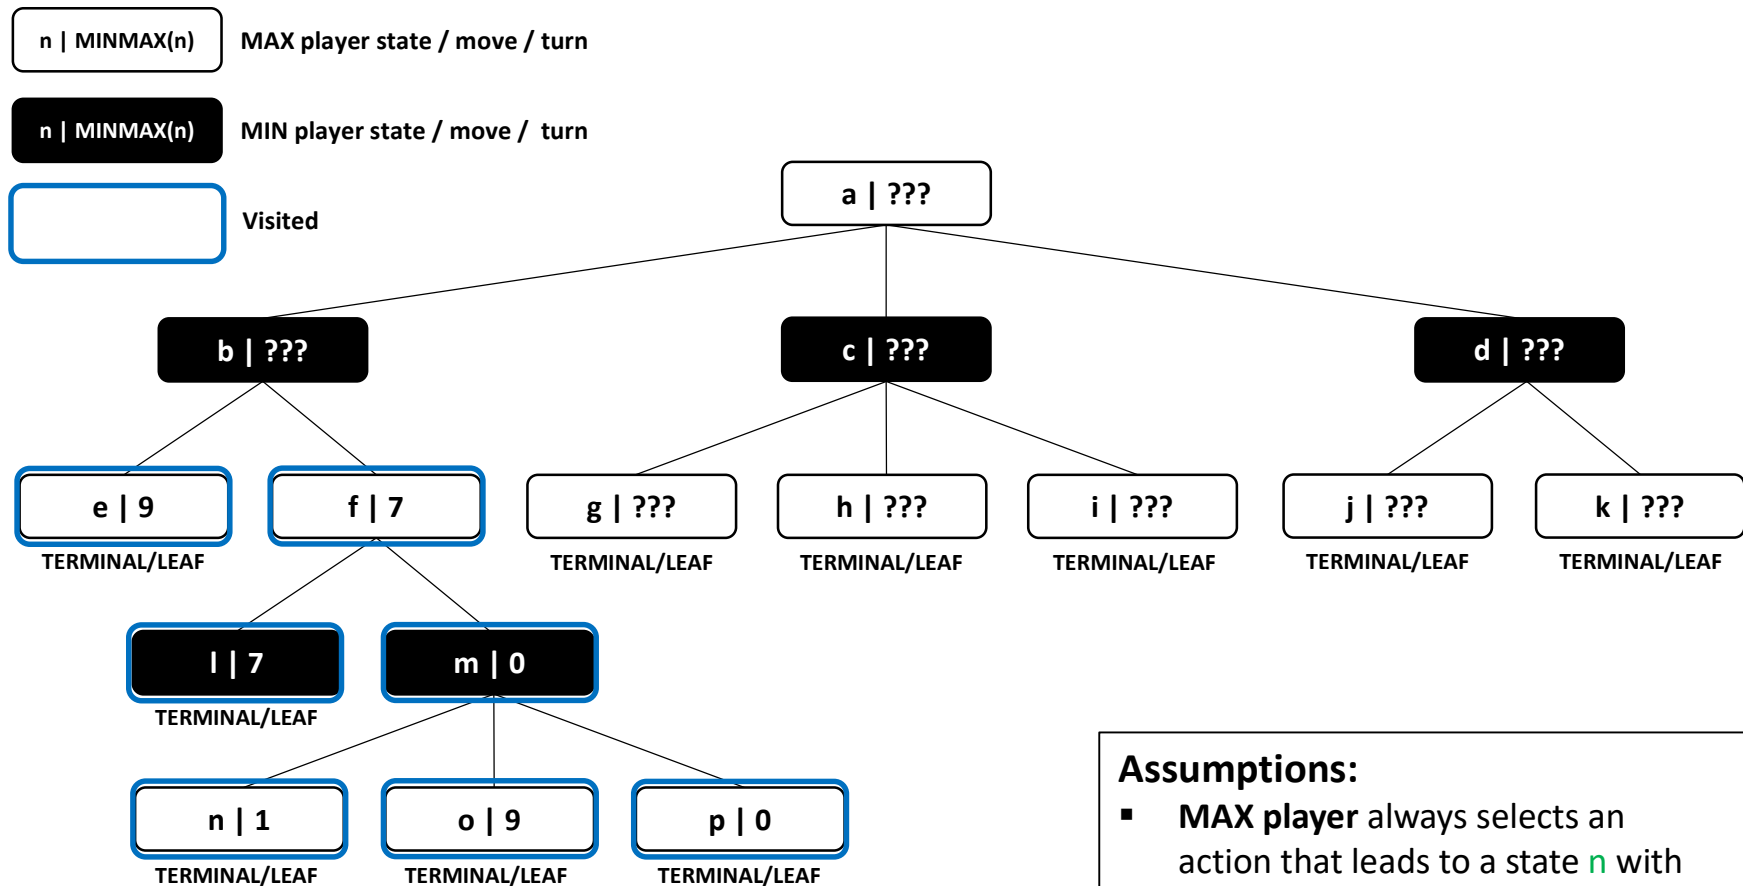
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- MAX player always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- MIN player always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- BOTH players always play optimally

Example MinMax Search Tree



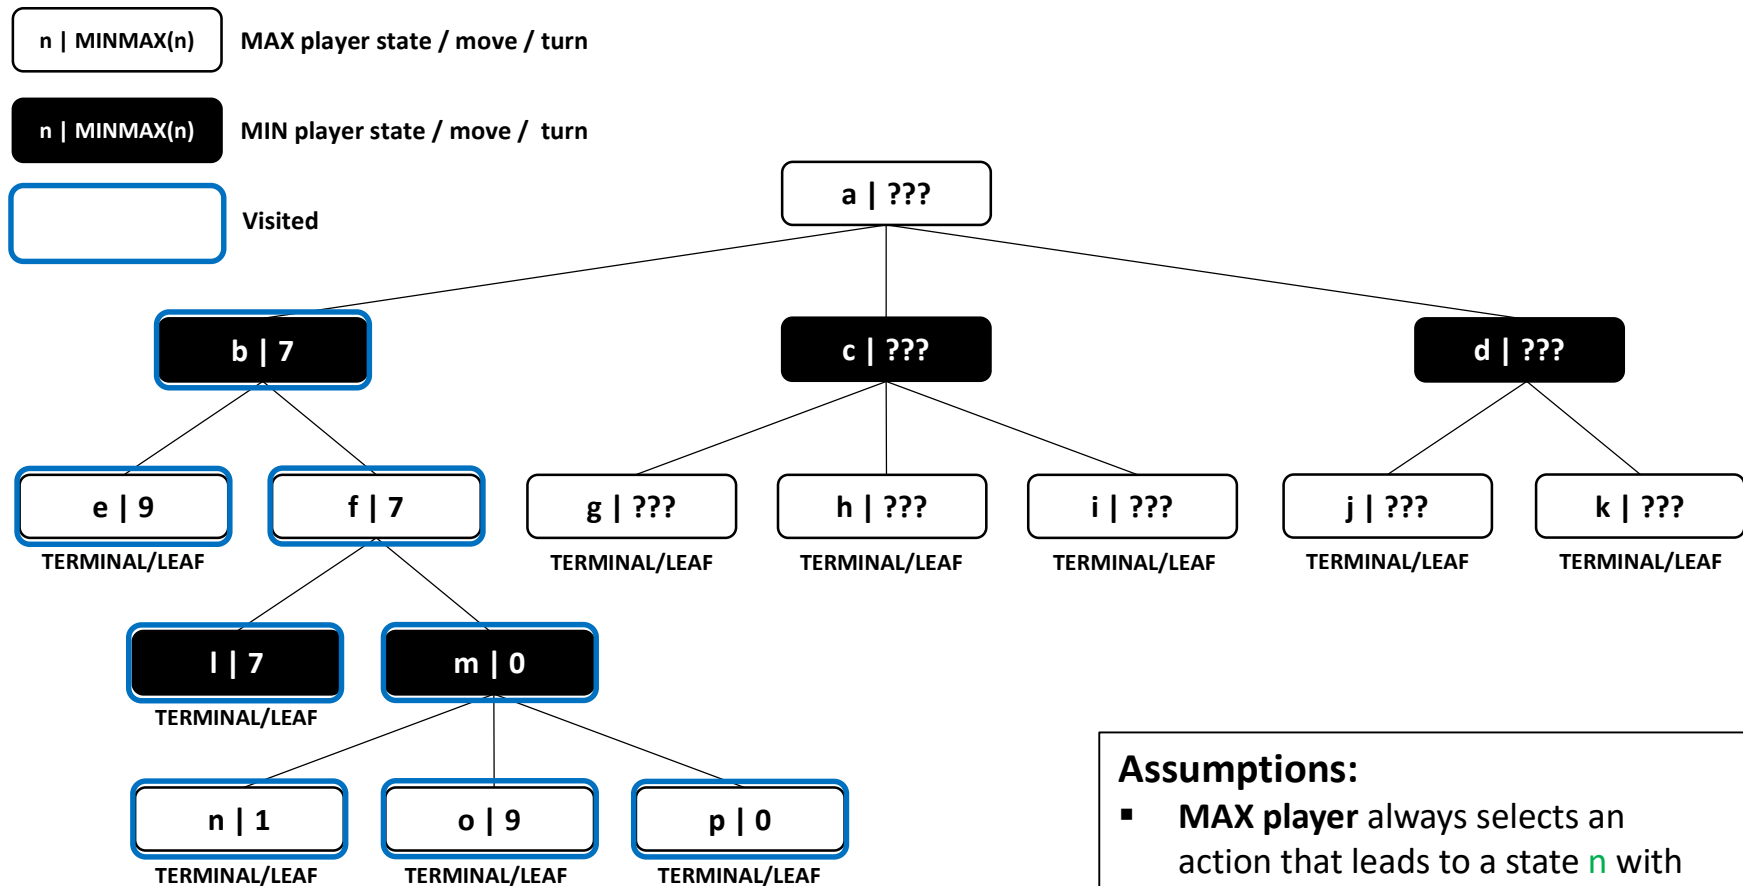
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINIMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINIMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



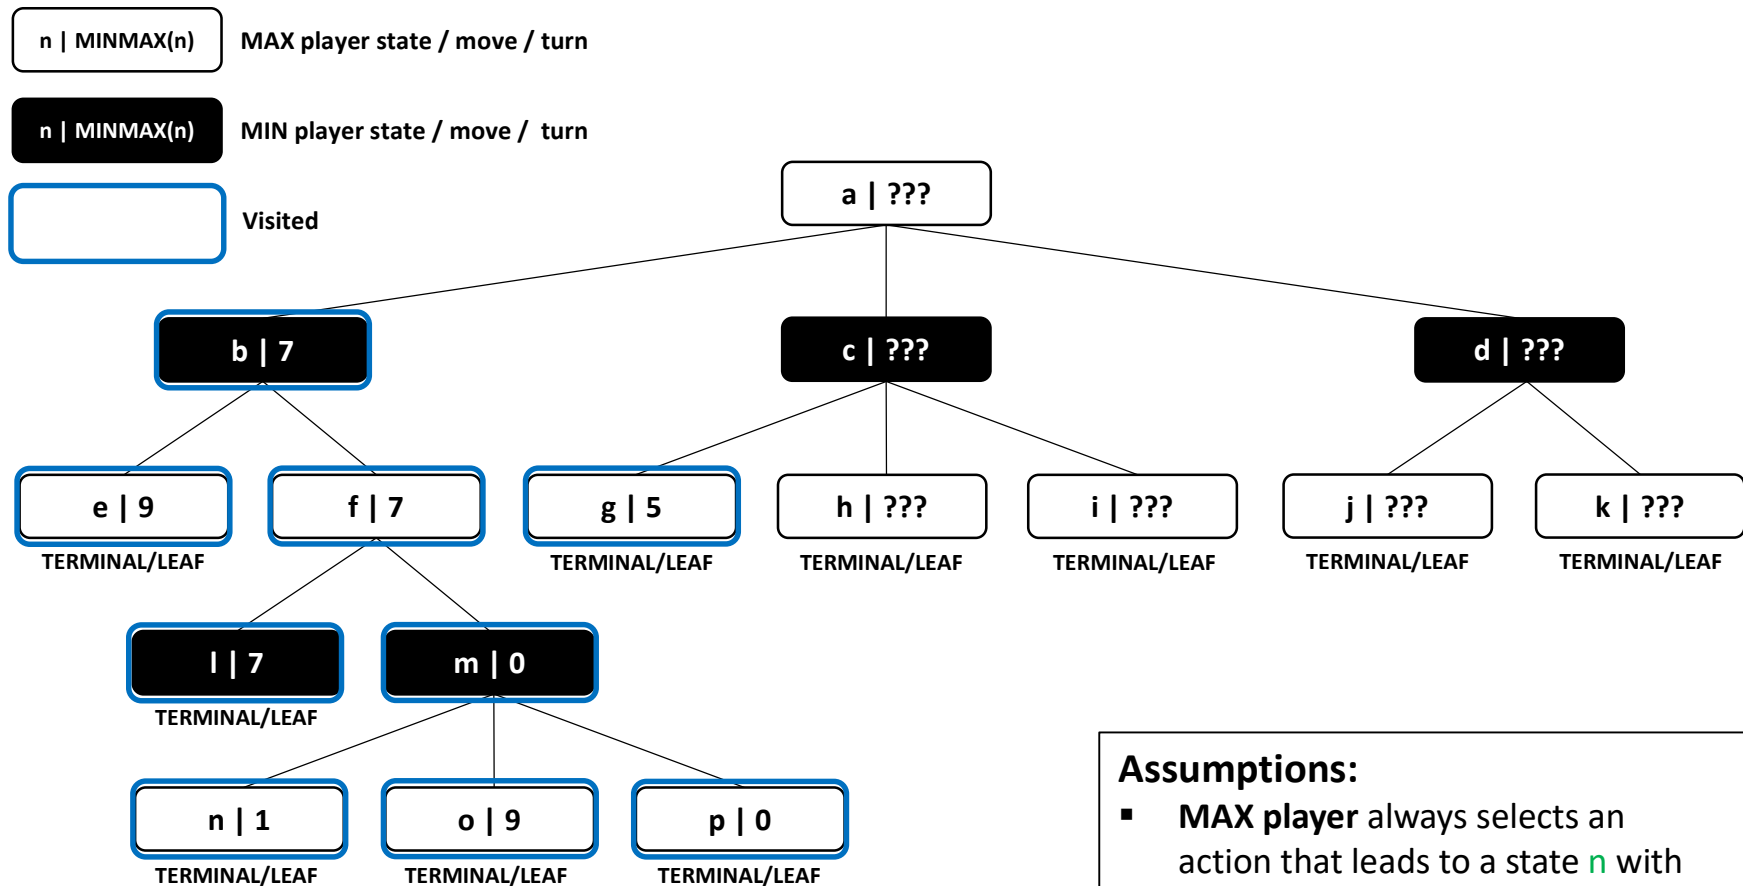
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINIMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINIMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



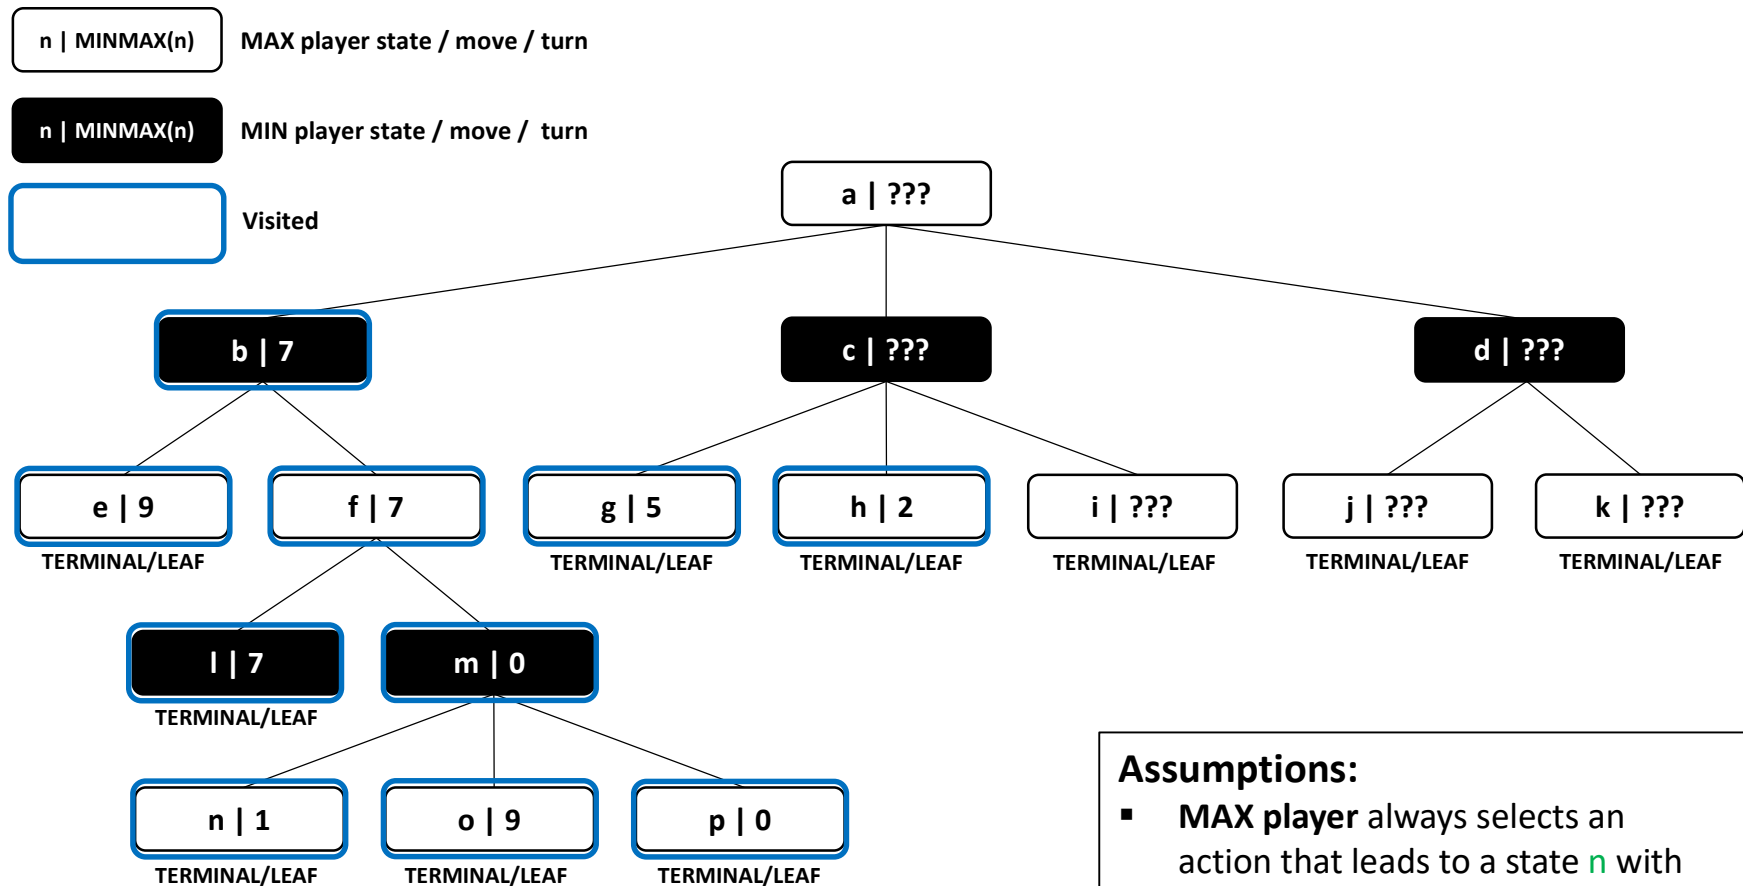
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINIMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINIMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



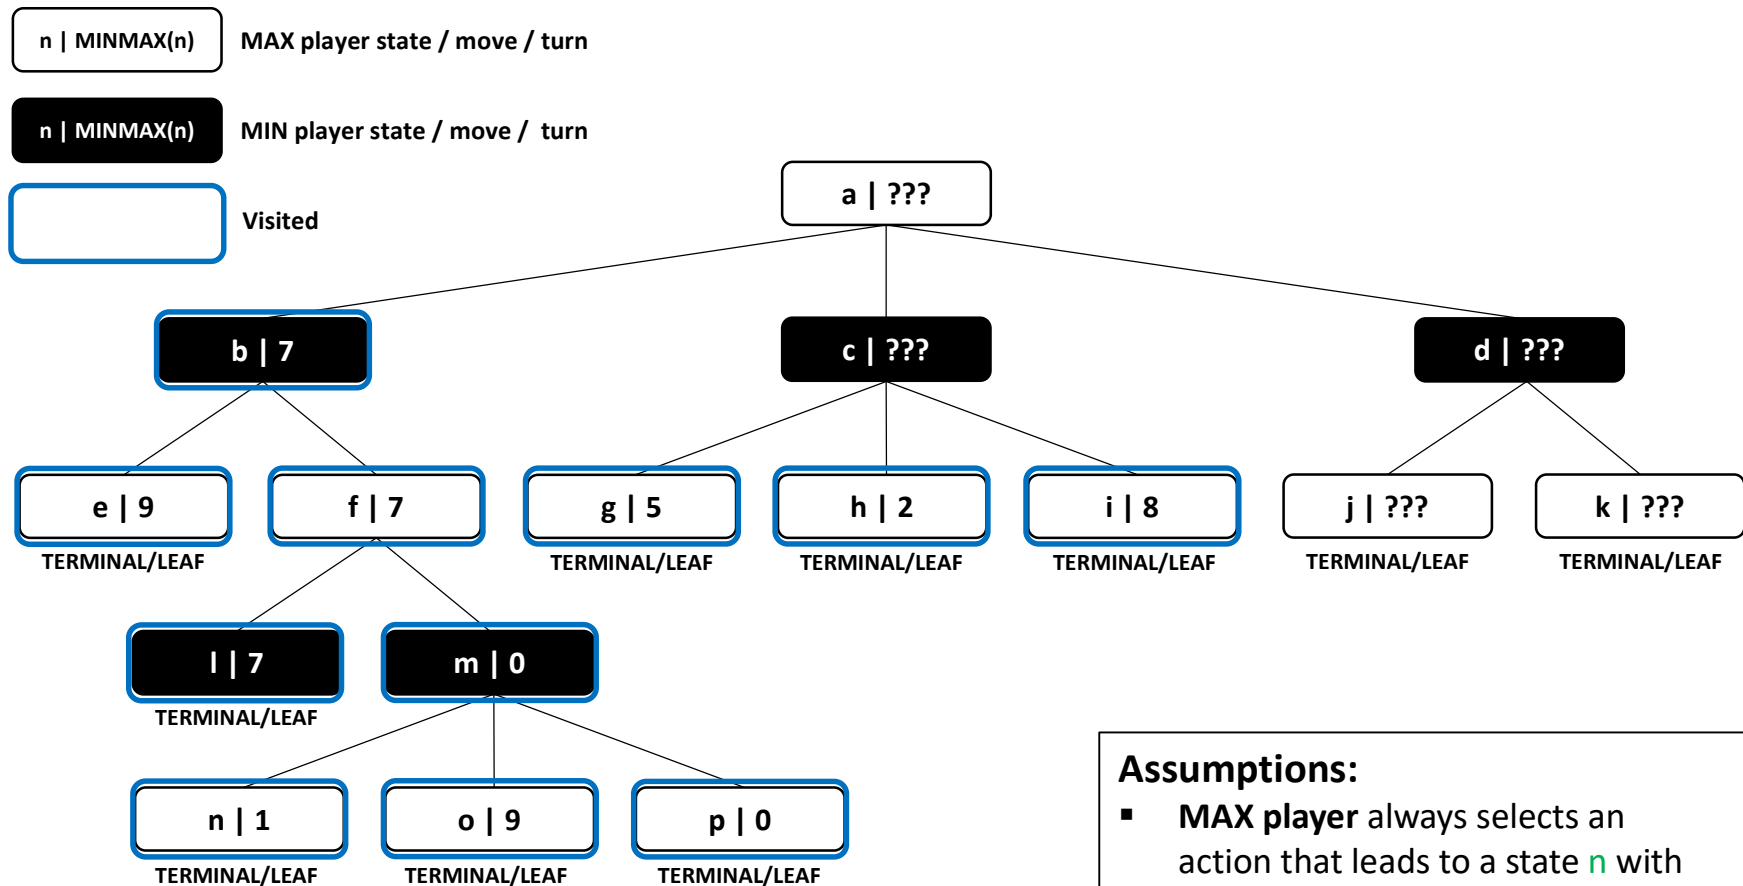
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINIMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINIMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



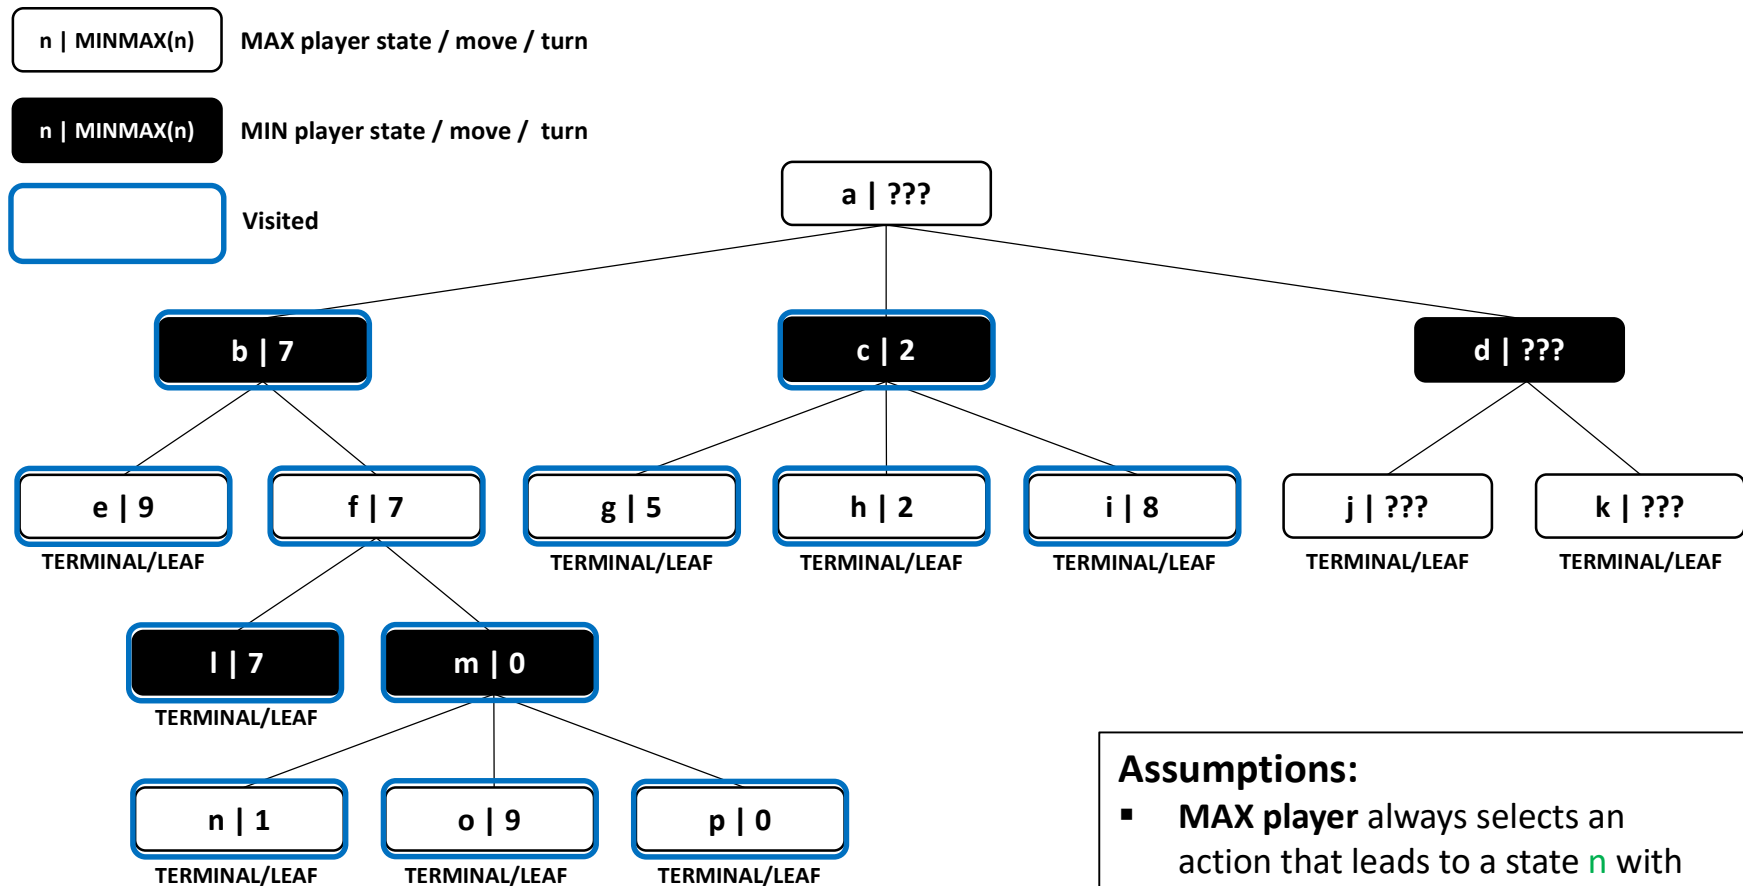
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- MAX player always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- MIN player always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- BOTH players always play optimally

Example MinMax Search Tree



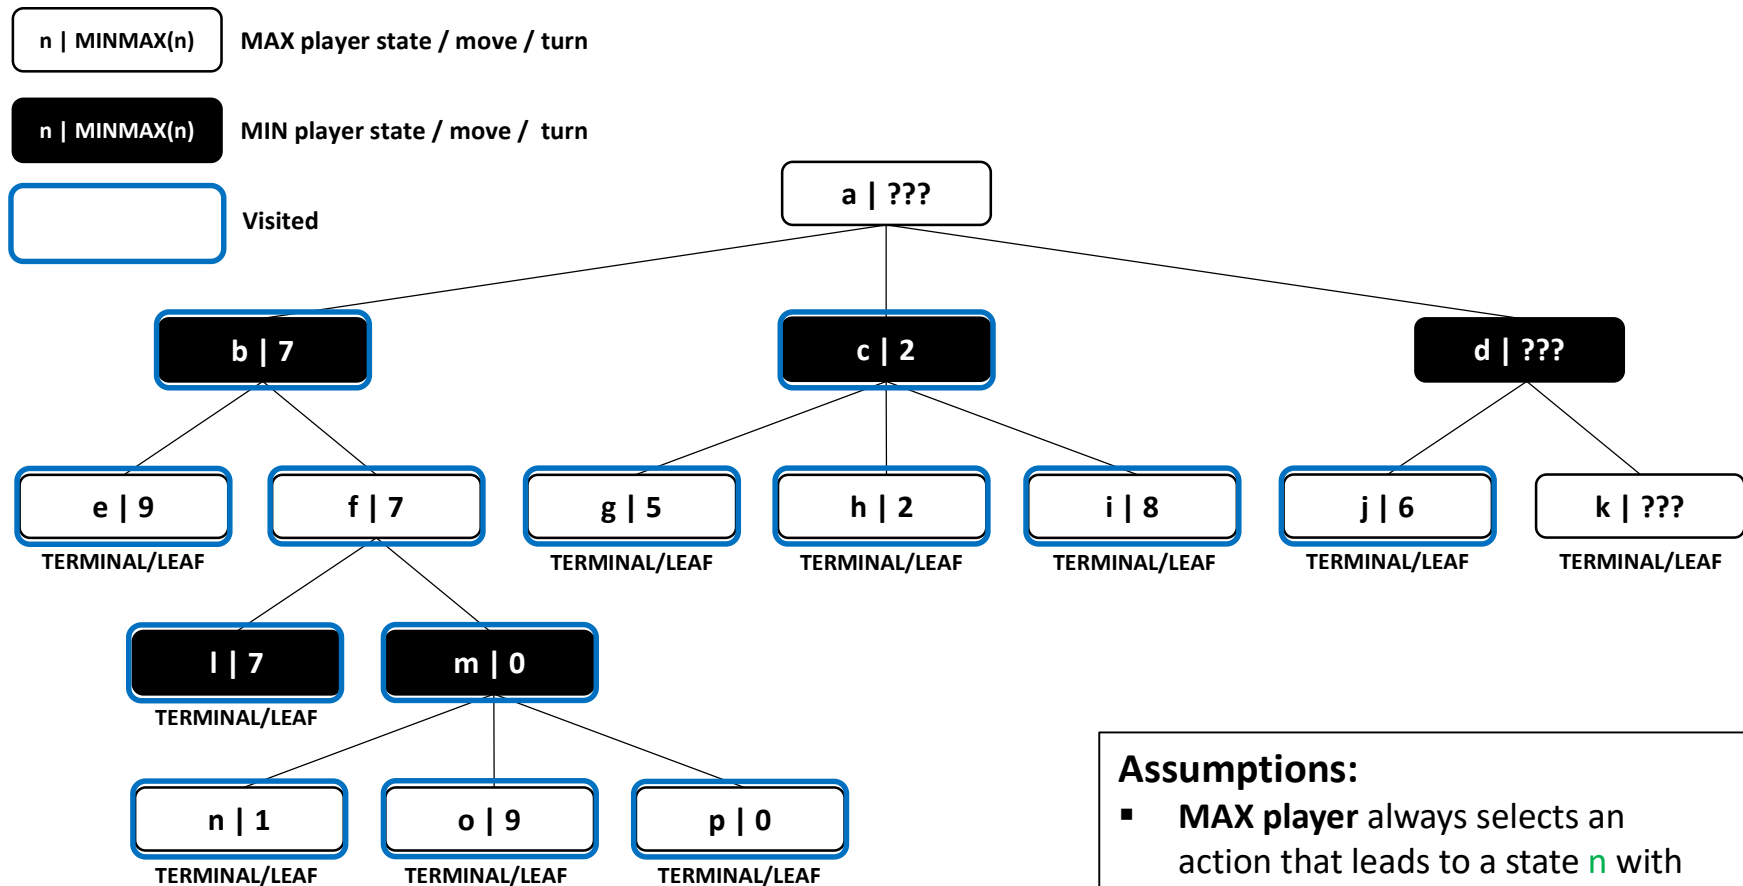
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



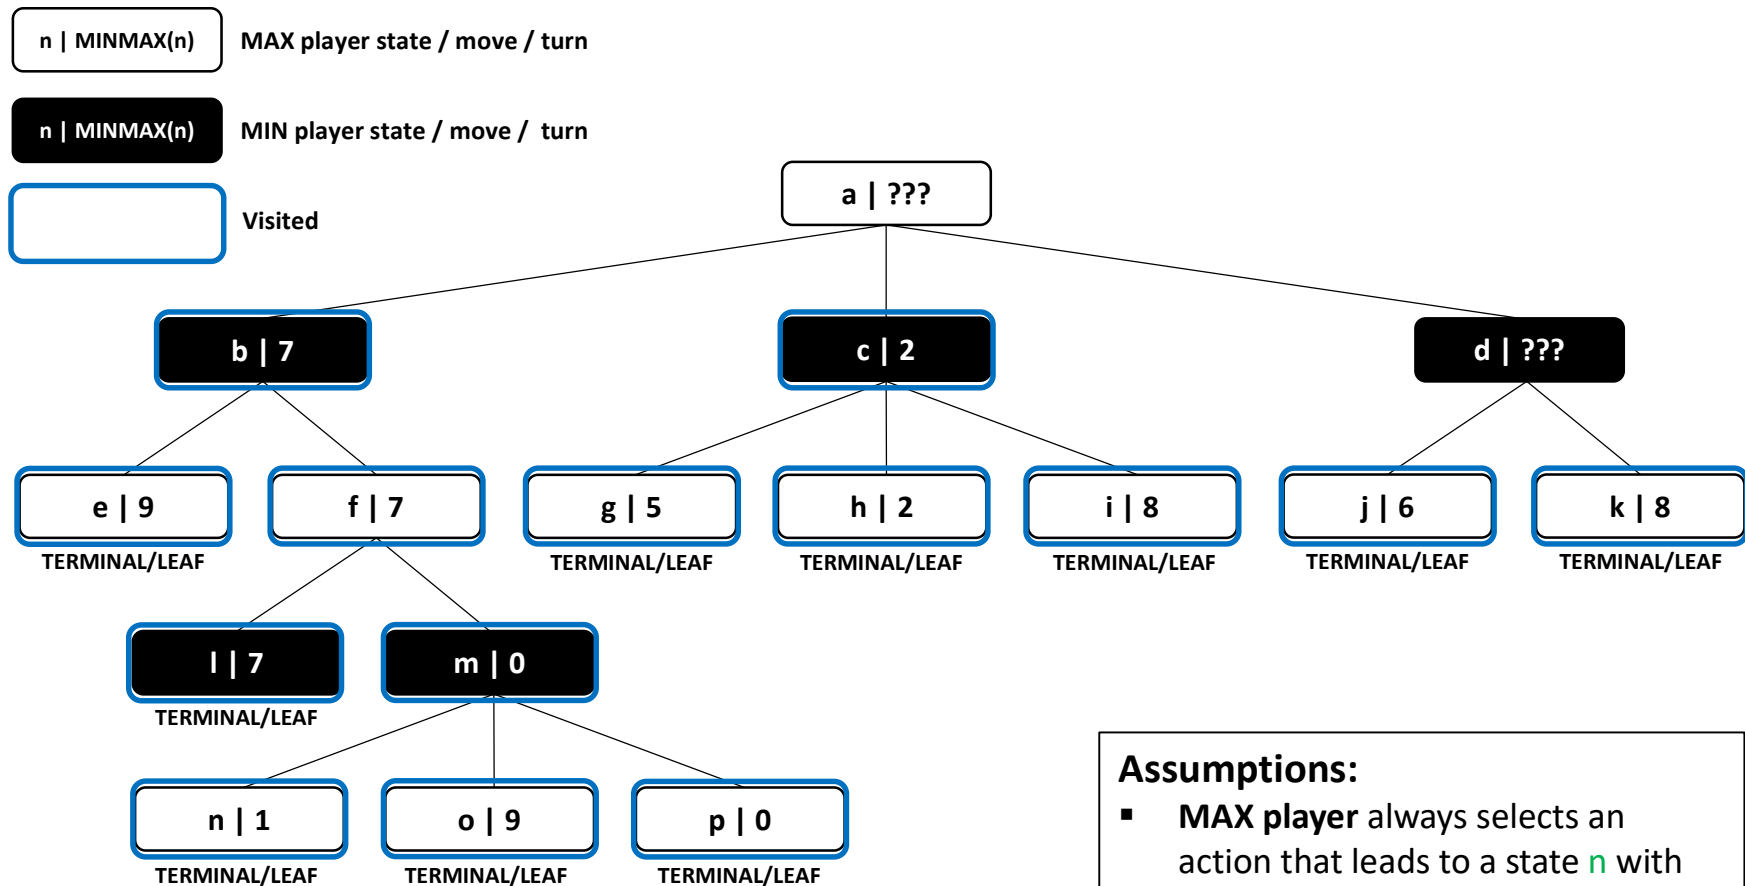
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- MAX player always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- MIN player always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- BOTH** players always play optimally

Example MinMax Search Tree



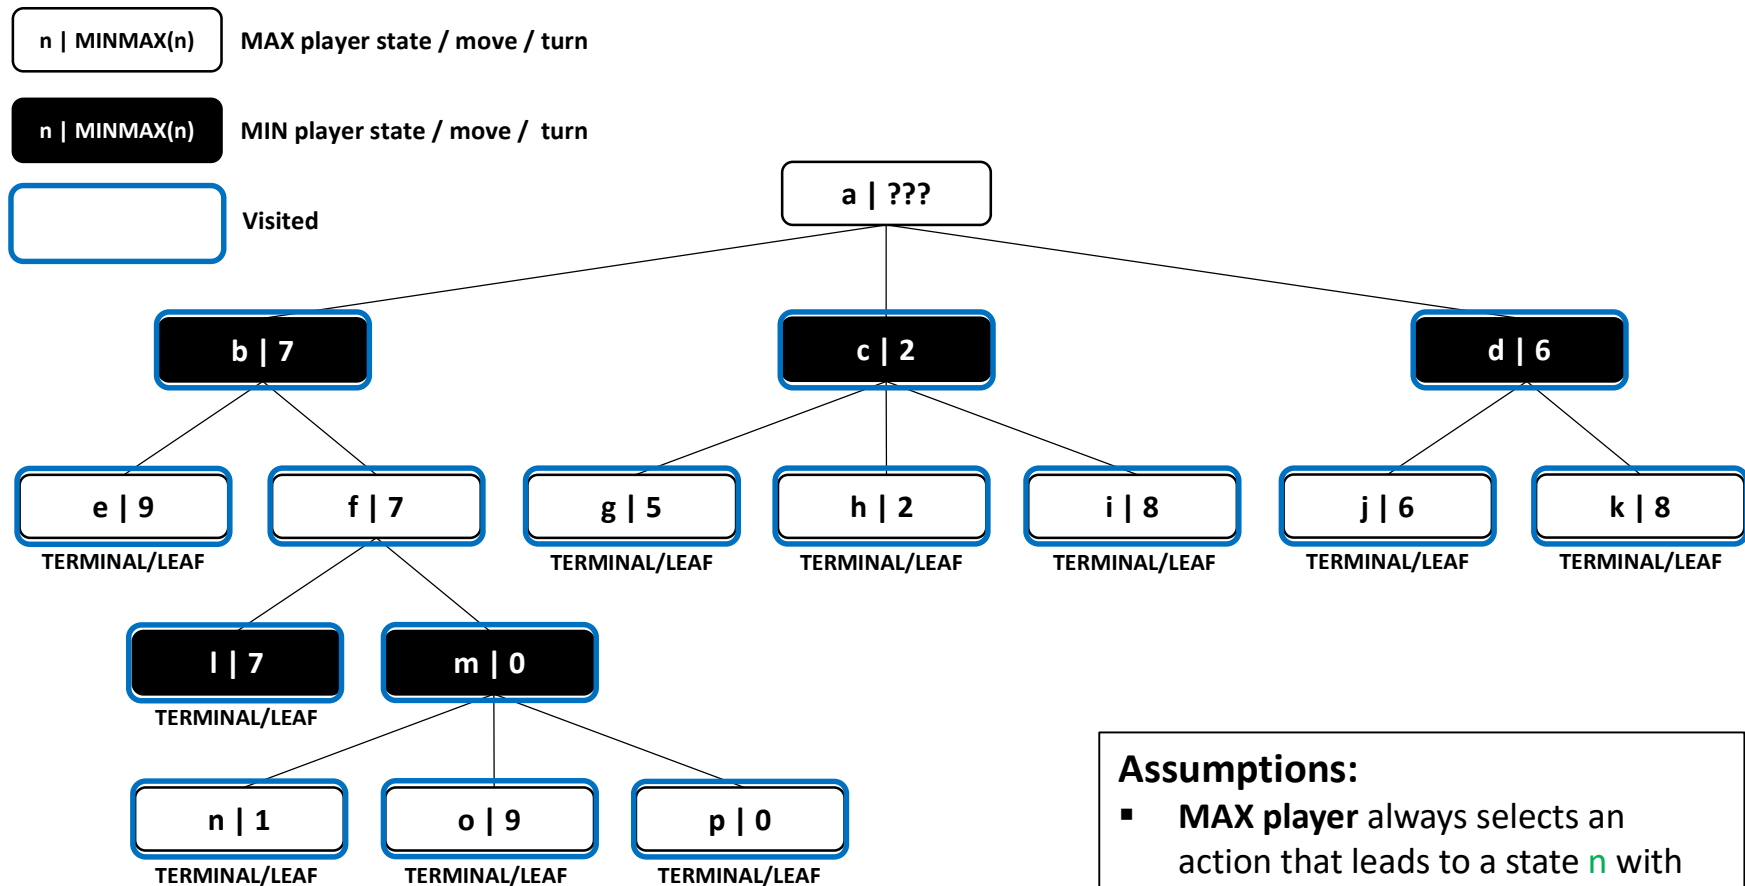
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



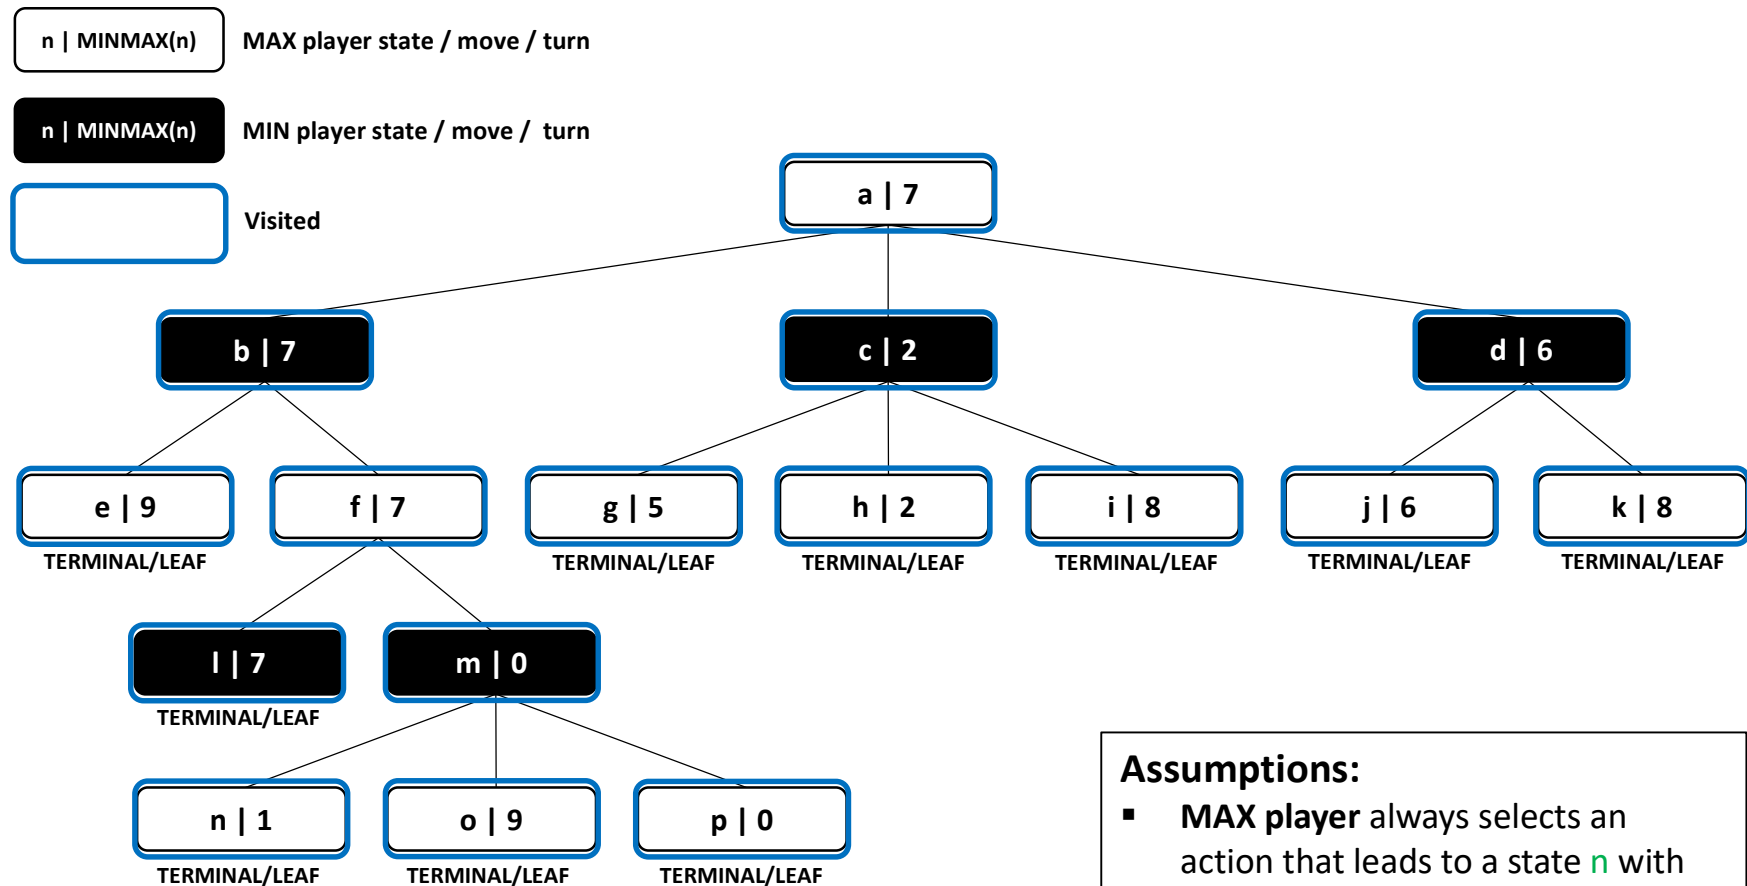
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- MAX player always selects an action that leads to a state n with **maximum** MINIMAX(n) value
- MIN player always selects an action that leads to a state n with **minimum** MINIMAX(n) value
- BOTH players always play optimally

Example MinMax Search Tree



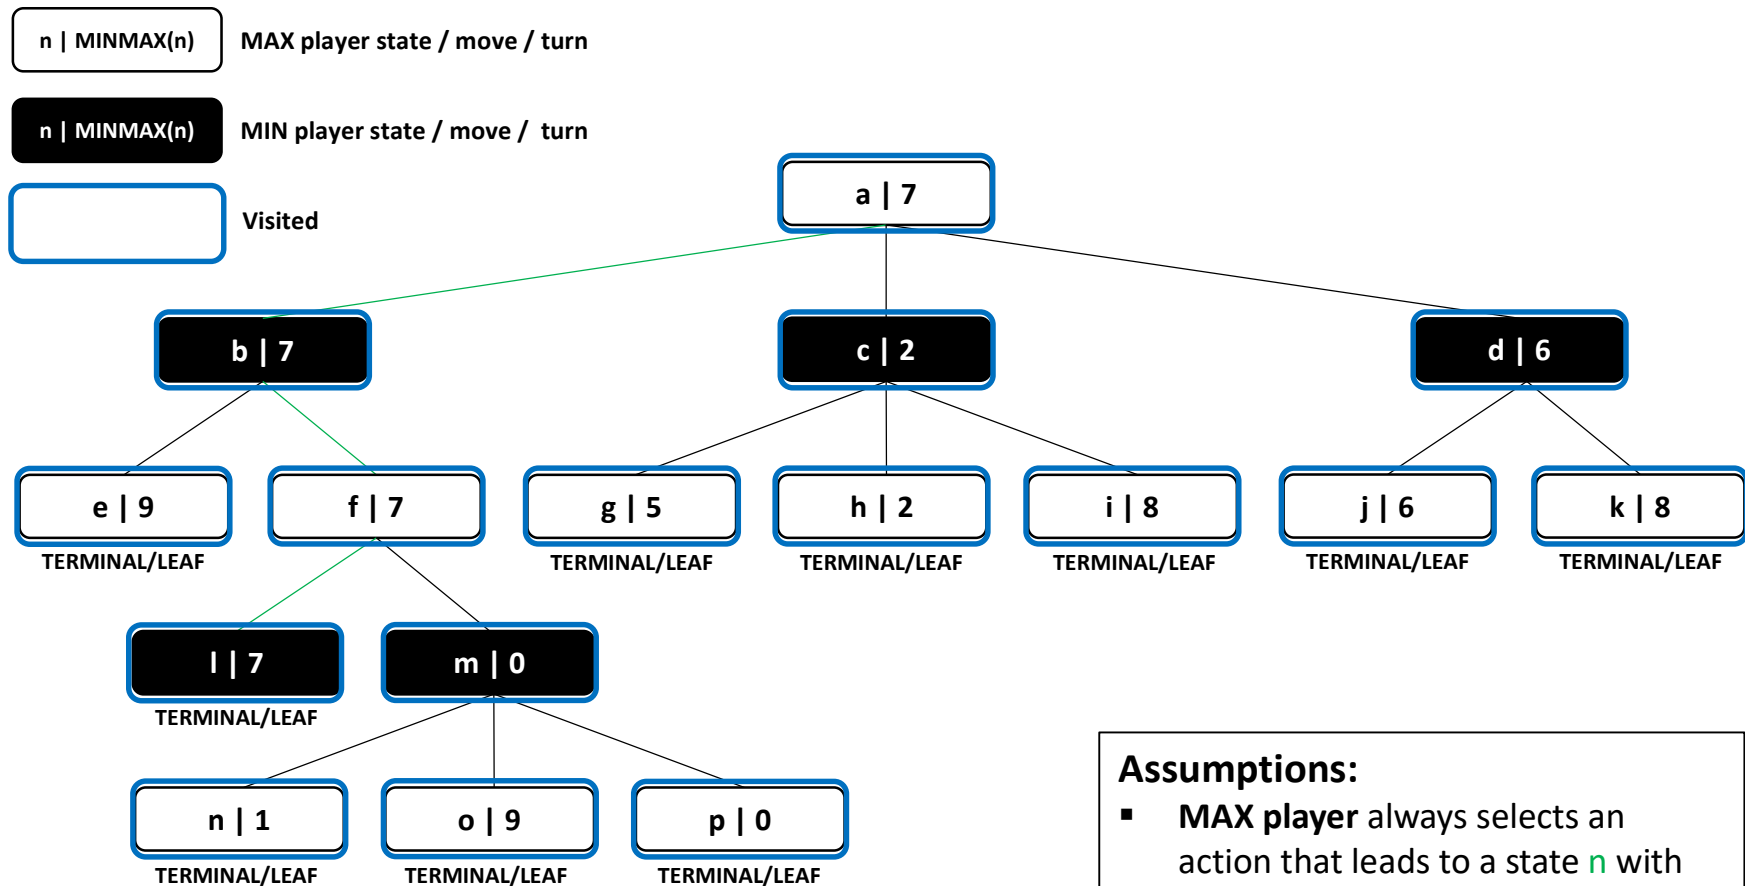
$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** $\text{MINMAX}(n)$ value
- **MIN player** always selects an action that leads to a state n with **minimum** $\text{MINMAX}(n)$ value
- **BOTH players** always play optimally

Example MinMax Search Tree



$$\text{MINMAX}(n) = \text{MINMAX}(\text{State}_n)$$

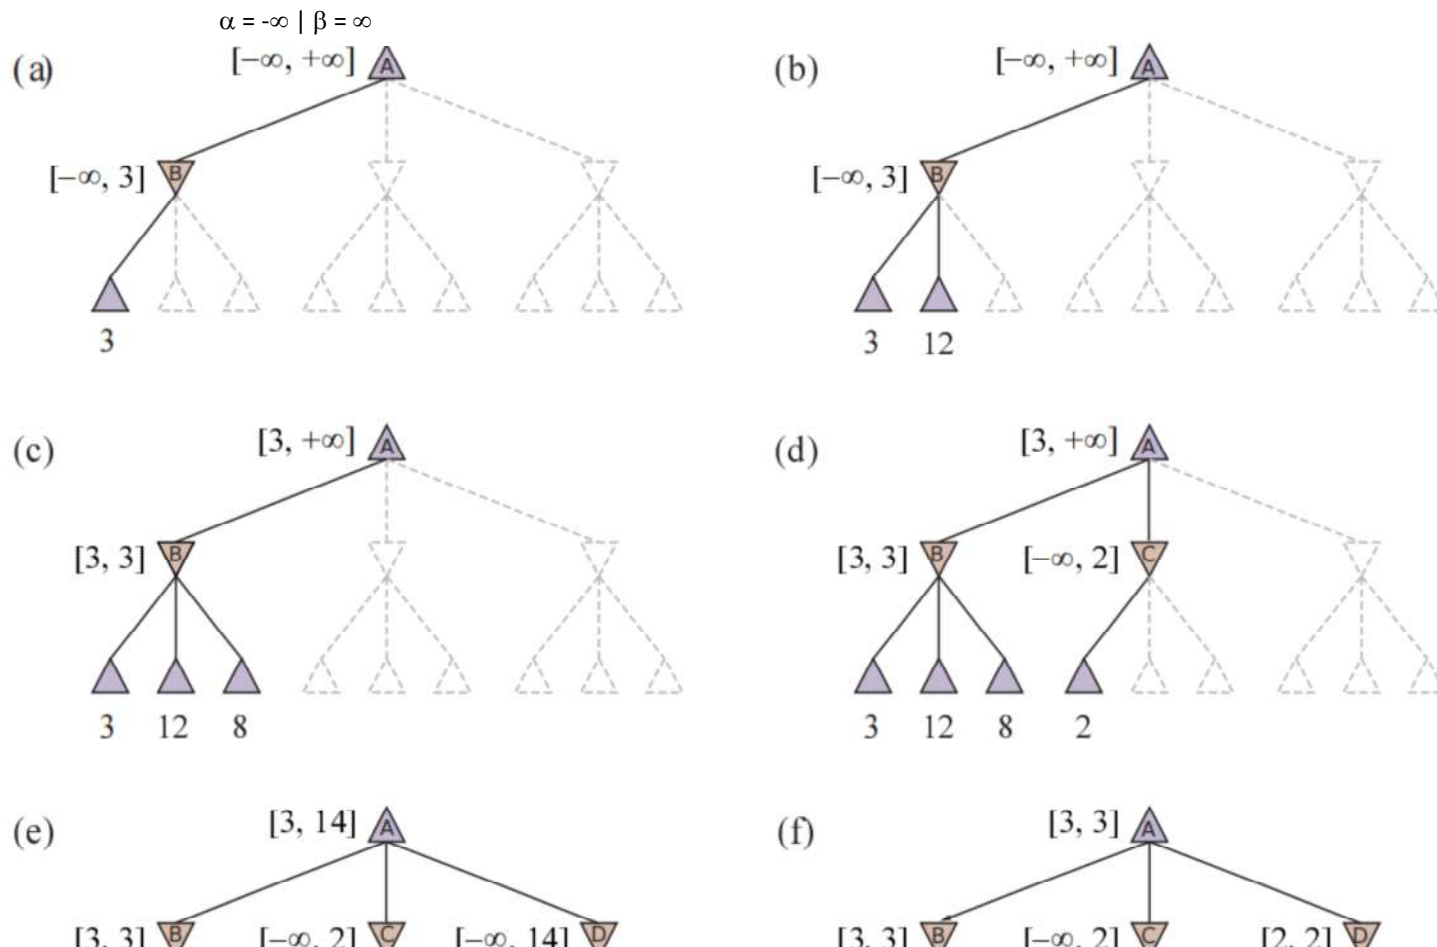
$$\text{MINMAX}(n) = \begin{cases} \text{UTILITY}(n, \text{MAX}), & \text{if } \text{ISTERMINAL}(n) \\ \max_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(n)} \text{MINMAX}(\text{RESULT}(n, a)), & \text{if } \text{TOMOVE}(s) = \text{MIN} \end{cases}$$

Assumptions:

- **MAX player** always selects an action that leads to a state n with **maximum** MINIMAX(n) value
- **MIN player** always selects an action that leads to a state n with **minimum** MINIMAX(n) value
- **BOTH players** always play optimally

MinMax: What is the Challenge?

Example MinMax with α - β Pruning



α : the value of the best (highest-value) choice we have found so far at any choice point along the path for MAX player ("at least")
 β : the value of the best (lowest-value) choice we have found so far at any choice point along the path for MIN player ("at most")

Example MinMax with α - β Pruning

